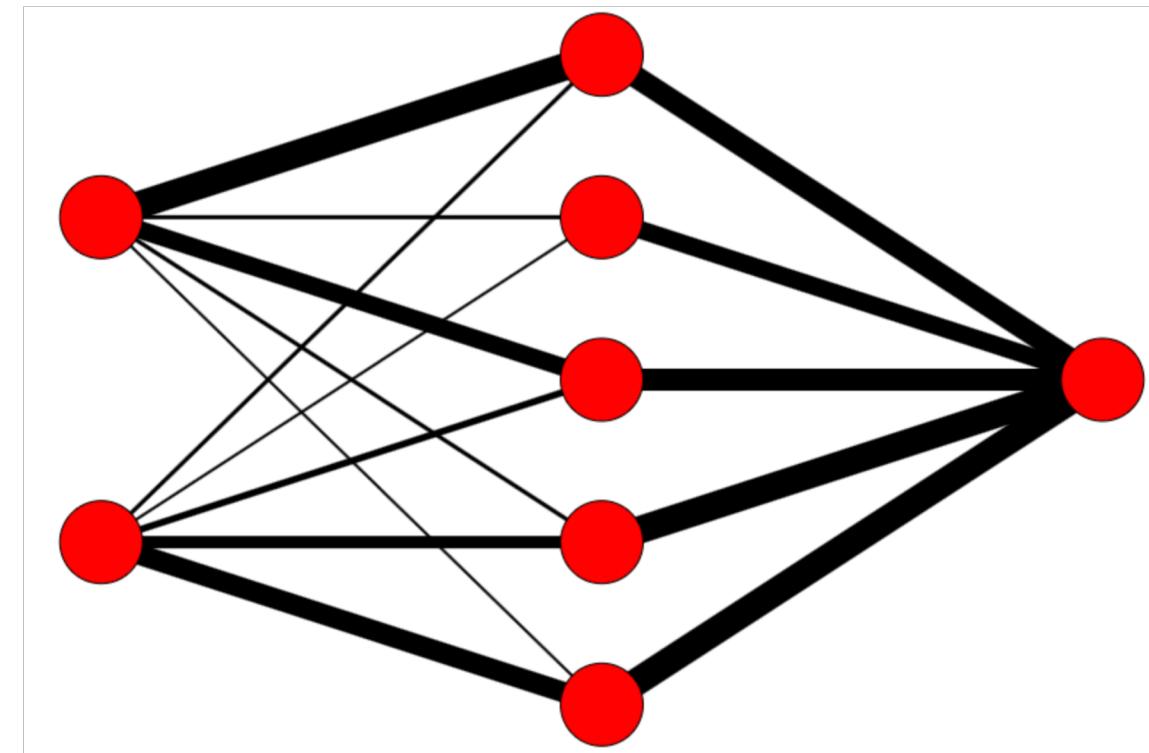
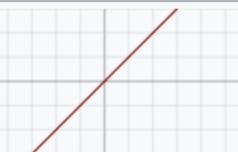
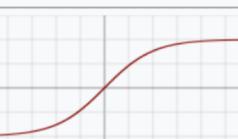
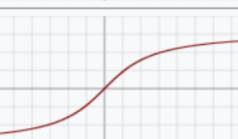


$$\text{output} = \text{Activation} \left( \sum (\text{weight} * \text{input}) + \text{bias} \right)$$

# Simple(st) network

<b>Input 1</b>	<b>Input 2</b>	<b>Target</b>
0	0	0
0	1	0
1	0	0
1	1	1

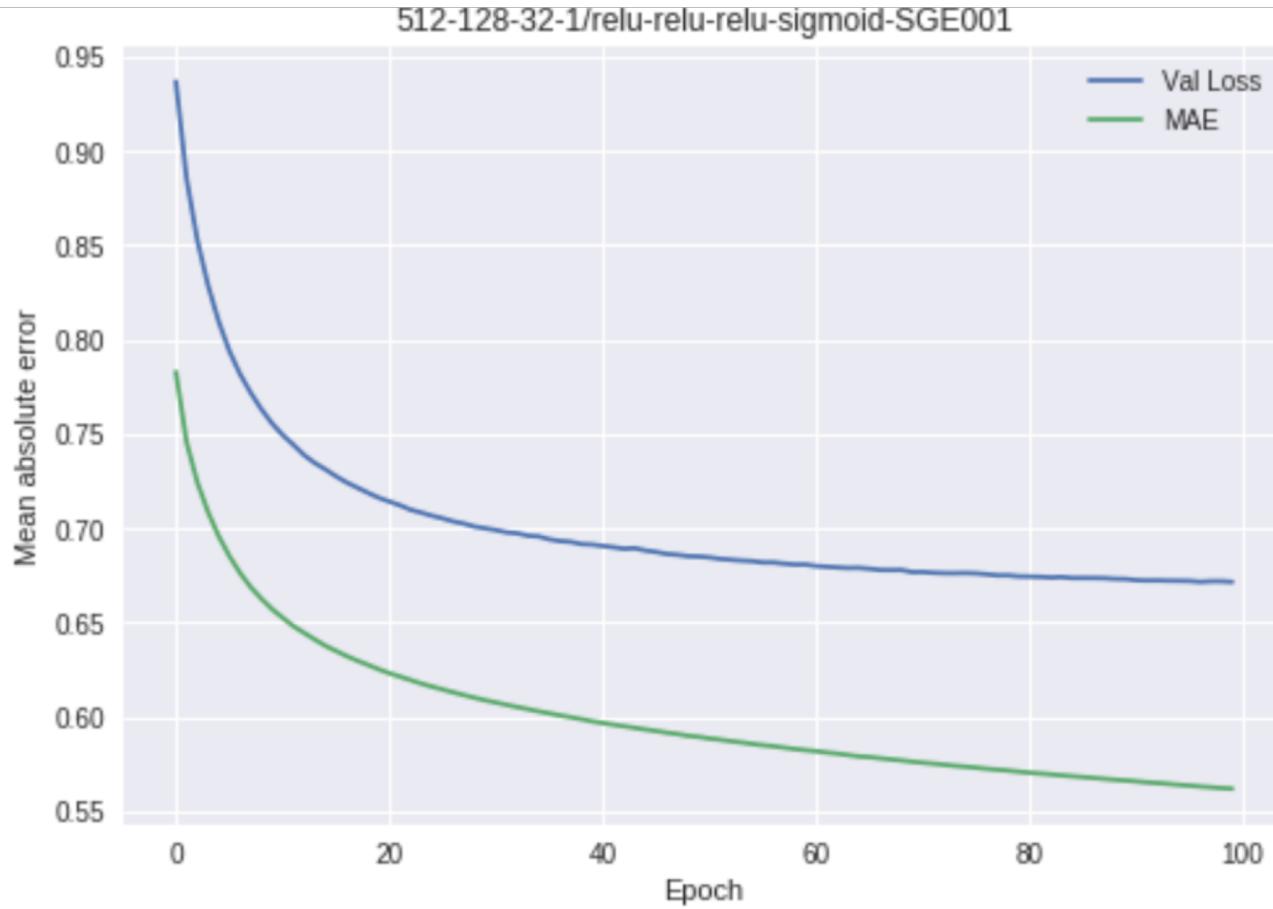


Name	Plot	Equation	Derivative (with respect to $x$ )	Range
Identity		$f(x) = x$	$f'(x) = 1$	$(-\infty, \infty)$
Binary step		$f(x) = \begin{cases} 0 & \text{for } x < 0 \\ 1 & \text{for } x \geq 0 \end{cases}$	$f'(x) = \begin{cases} 0 & \text{for } x \neq 0 \\ ? & \text{for } x = 0 \end{cases}$	$\{0, 1\}$
Logistic (a.k.a. Sigmoid or Soft step)		$f(x) = \sigma(x) = \frac{1}{1 + e^{-x}}$ [1]	$f'(x) = f(x)(1 - f(x))$	$(0, 1)$
TanH		$f(x) = \tanh(x) = \frac{(e^x - e^{-x})}{(e^x + e^{-x})}$	$f'(x) = 1 - f(x)^2$	$(-1, 1)$
ArcTan		$f(x) = \tan^{-1}(x)$	$f'(x) = \frac{1}{x^2 + 1}$	$\left(-\frac{\pi}{2}, \frac{\pi}{2}\right)$
Rectified linear unit (ReLU) <sup>[10]</sup>		$f(x) = \begin{cases} 0 & \text{for } x < 0 \\ x & \text{for } x \geq 0 \end{cases}$	$f'(x) = \begin{cases} 0 & \text{for } x < 0 \\ 1 & \text{for } x \geq 0 \end{cases}$	$[0, \infty)$
Leaky rectified linear unit (Leaky ReLU) <sup>[11]</sup>		$f(x) = \begin{cases} 0.01x & \text{for } x < 0 \\ x & \text{for } x \geq 0 \end{cases}$	$f'(x) = \begin{cases} 0.01 & \text{for } x < 0 \\ 1 & \text{for } x \geq 0 \end{cases}$	$(-\infty, \infty)$

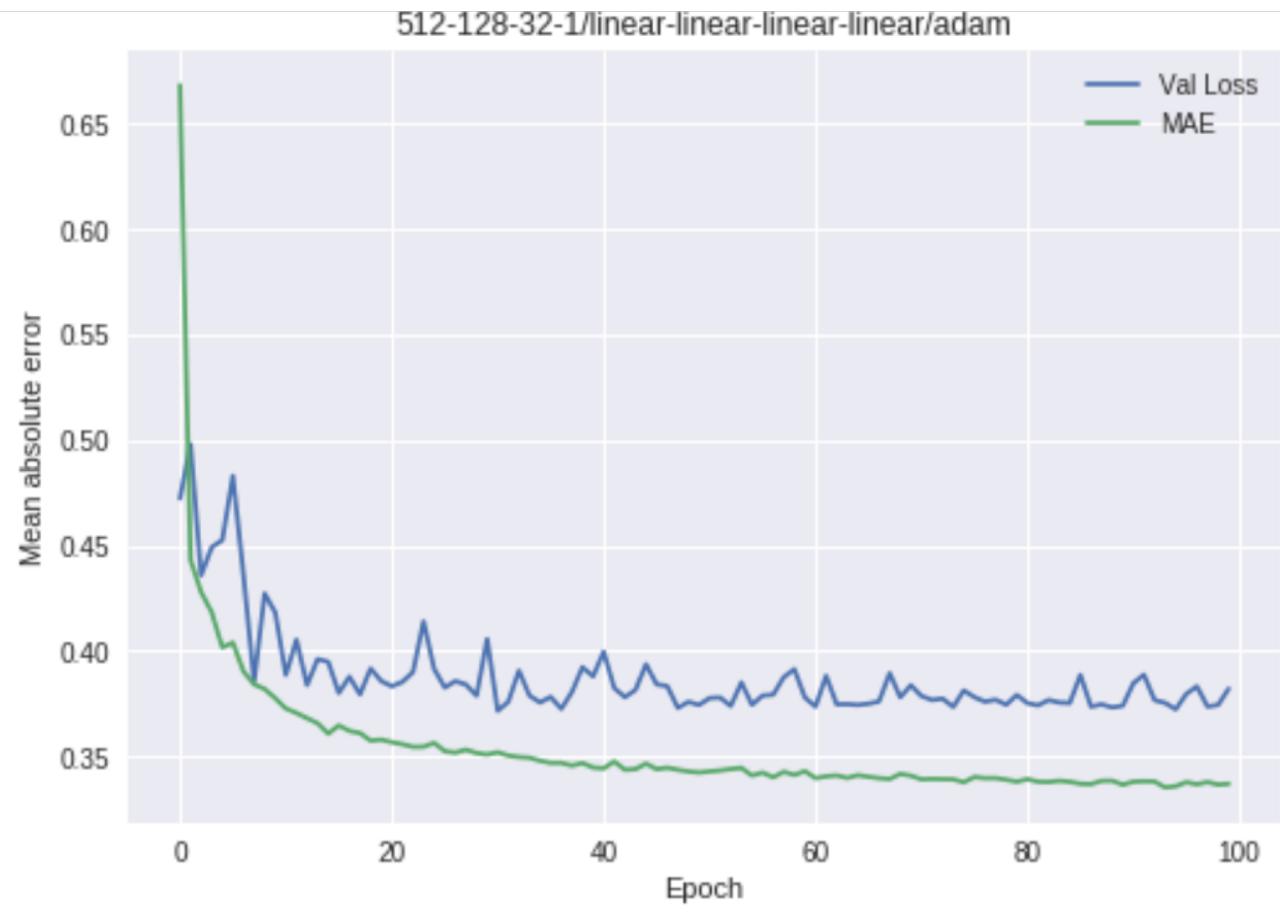
## Example code and data

- <https://bit.ly/2RiQesC>
- <https://colab.research.google.com>

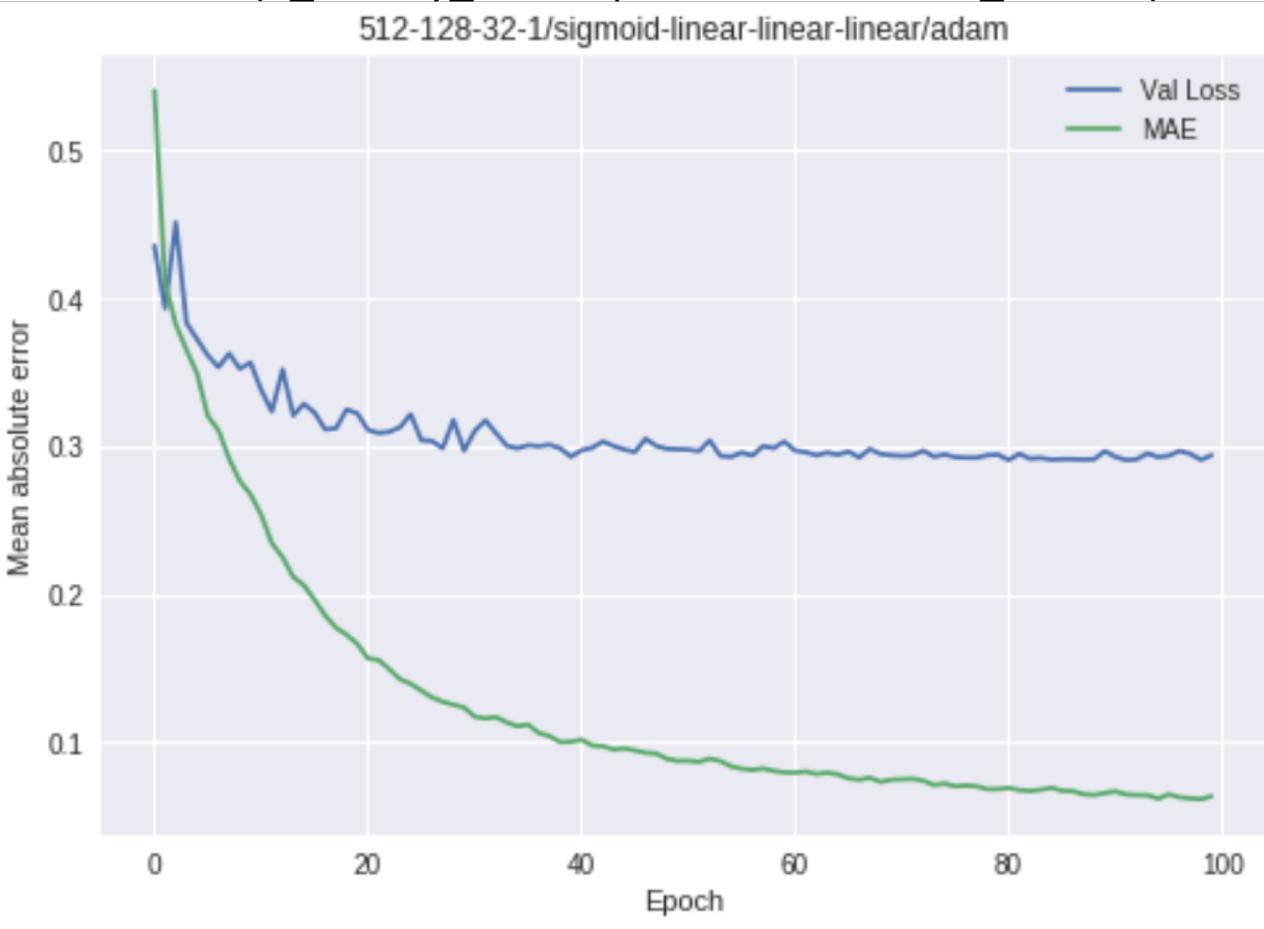
```
model = Sequential()  
model.add(Dense(units=512, activation='relu', input_dim=num_columns-1))  
model.add(Dense(units=128, activation='relu'))  
model.add(Dense(units=32, activation='relu'))  
model.add(Dense(units=1, activation='sigmoid'))  
  
model.compile(loss='mae', optimizer=SGD(lr=0.001), metrics=["mae"])  
model.fit(X_train, y_train, epochs=5000, batch_size=32)
```



```
model = Sequential()  
model.add(Dense(units=512, activation='linear', input_dim=num_columns-1))  
model.add(Dense(units=128, activation='linear'))  
model.add(Dense(units=32, activation='linear'))  
model.add(Dense(units=1, activation='linear'))  
  
model.compile(loss='mae', optimizer='adam', metrics=["mse","mae"])  
model.fit(X_train, y_train, epochs=5000, batch_size=32)
```



```
model = Sequential()  
model.add(Dense(units=512, activation='sigmoid', input_dim=num_columns-1))  
model.add(Dense(units=128, activation='linear'))  
model.add(Dense(units=32, activation='linear'))  
model.add(Dense(units=1, activation='linear'))  
  
model.compile(loss='mae', optimizer='adam', metrics=["mse","mae"])  
model.fit(X_train, y_train, epochs=5000, batch_size=32)
```



## # Learning Boolean logic

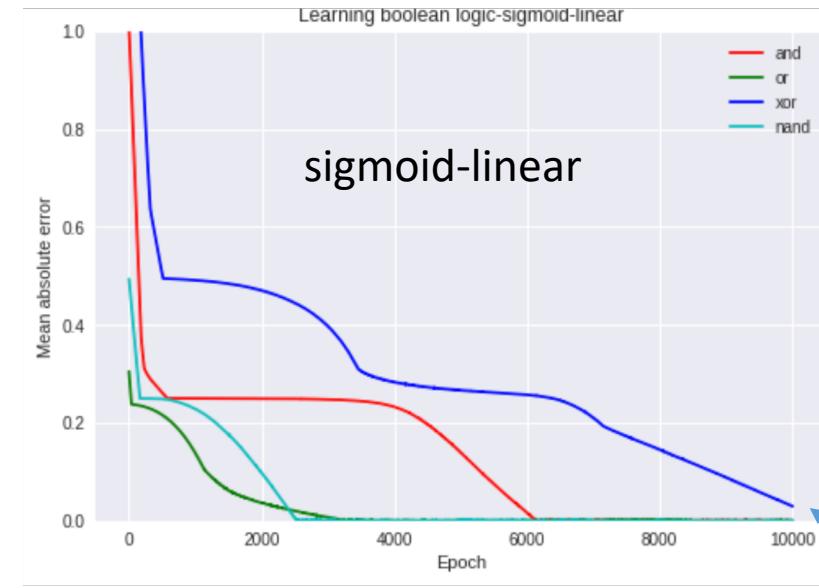
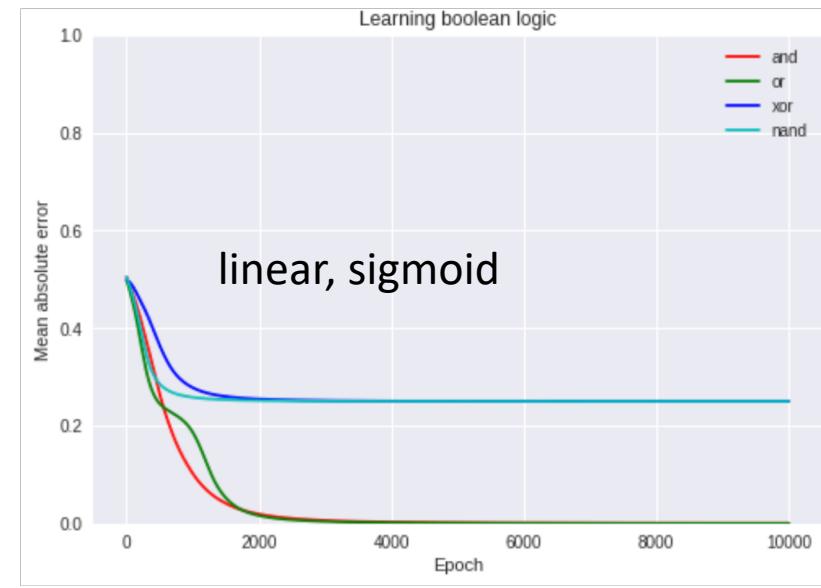
```
act1 = 'sigmoid'  
act2 = 'linear'  
optimizer = 'adam'  
epochs = 10000  
batch_size = 32
```

```
boolX = numpy.array([[0,1],[1,1],[1,0],[0,0]])  
andy = numpy.array([0,1,0,0])  
ory = numpy.array([1,1,1,0])  
xory = numpy.array([1,0,1,0])  
nandy = numpy.array([1,0,1,1])
```

```
model= Sequential()  
model.add(Dense(units=3, activation=act1, input_dim=2))  
model.add(Dense(units=1, activation=act2))
```

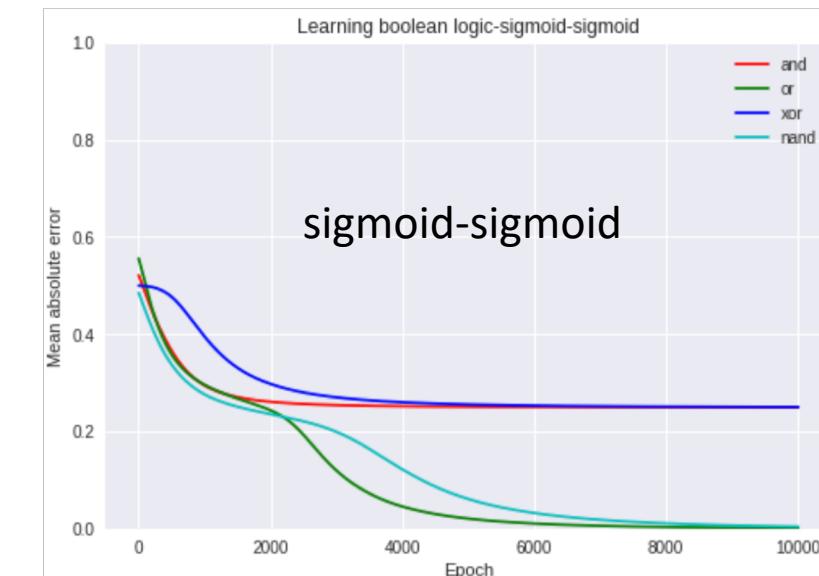
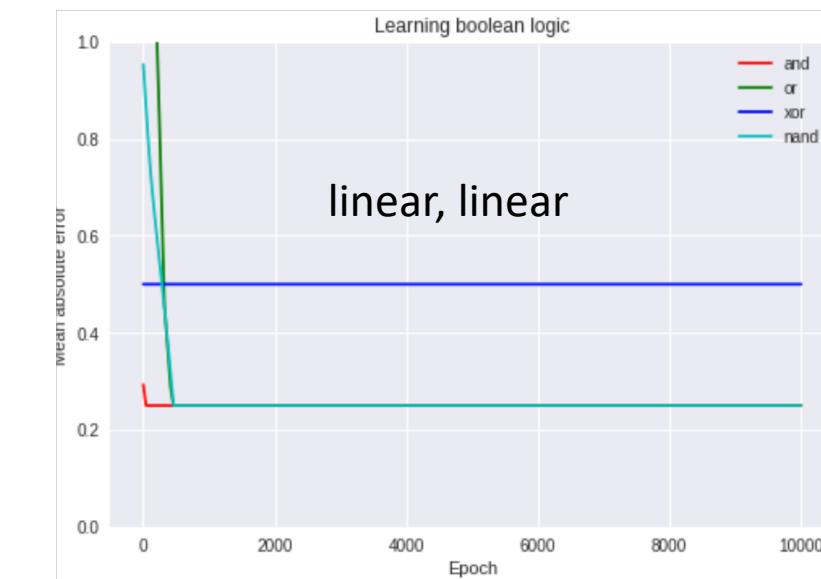
```
model.compile(loss='mae', optimizer=optimizer, metrics=["mse", "mae"])  
h1= boolmod1.fit(boolX, andy, epochs=epochs, batch_size=batch_size, verbose=0)  
h2= boolmod1.fit(boolX, ory, epochs=epochs, batch_size=batch_size, verbose=0)  
h3= boolmod1.fit(boolX, xory, epochs=epochs, batch_size=batch_size, verbose=0)  
h4= boolmod1.fit(boolX, nandy, epochs=epochs, batch_size=batch_size, verbose=0)
```

# Activation functions matter

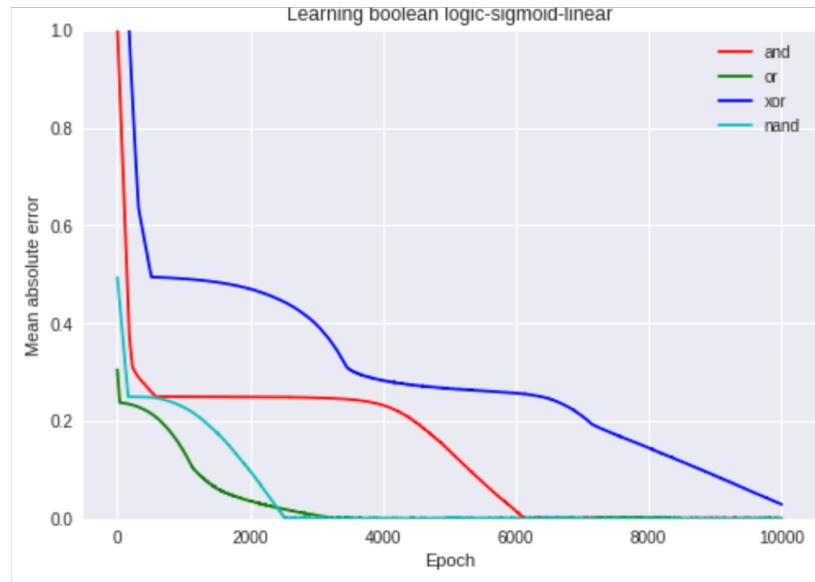


All with:  
batch-size = 32  
optimizer=adam

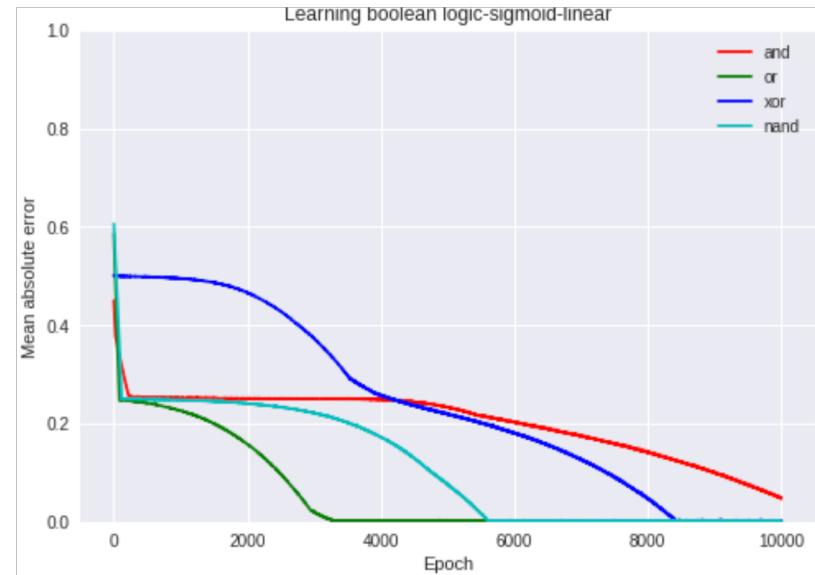
It gets there at around 2,000 epochs



# Batch size matters



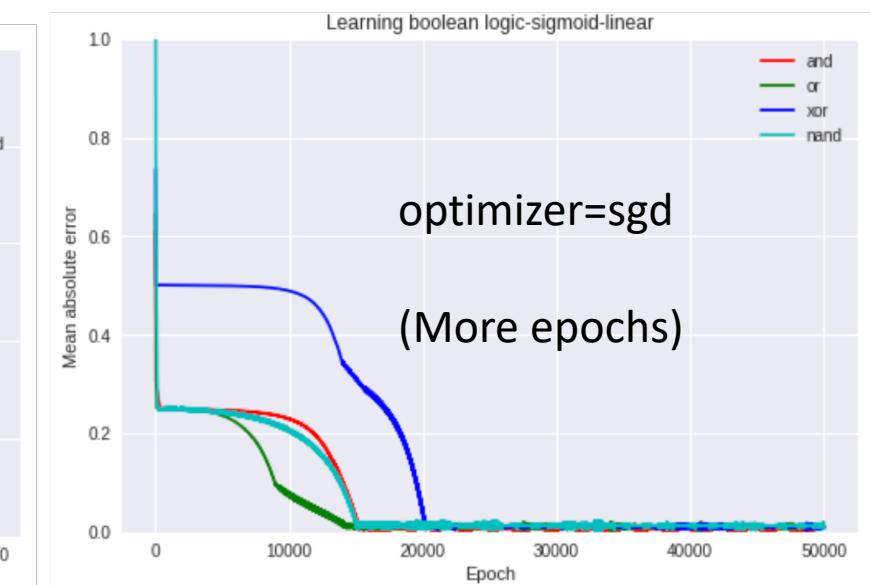
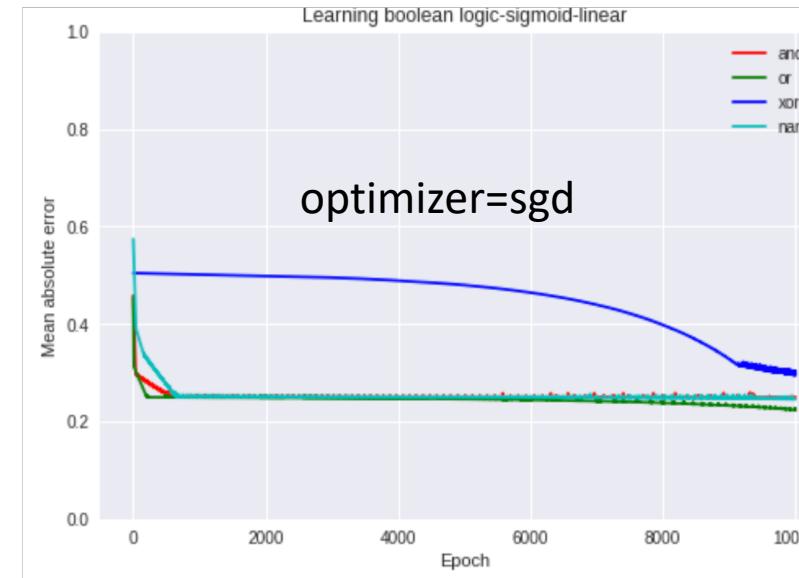
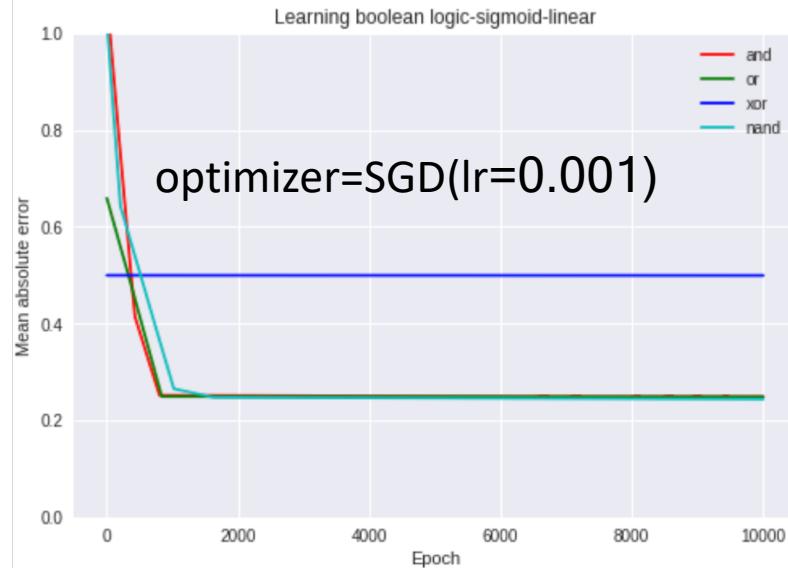
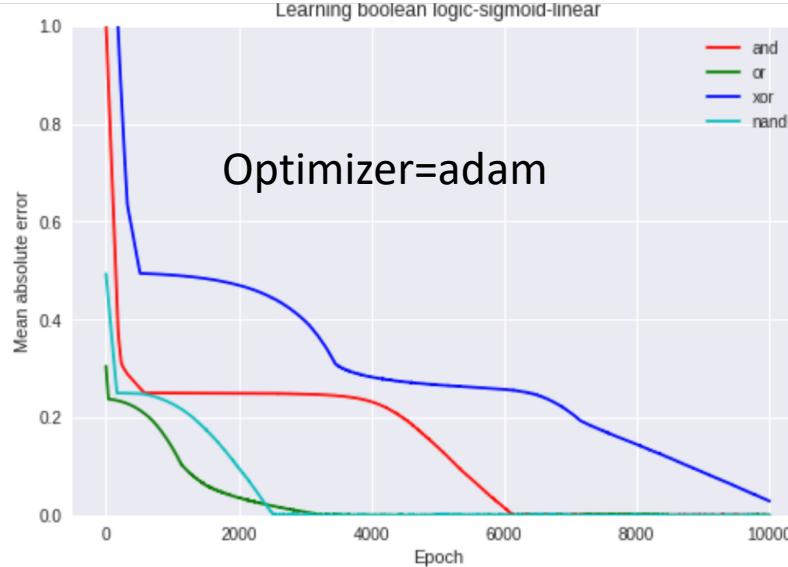
batch-size = 32



batch-size = 1

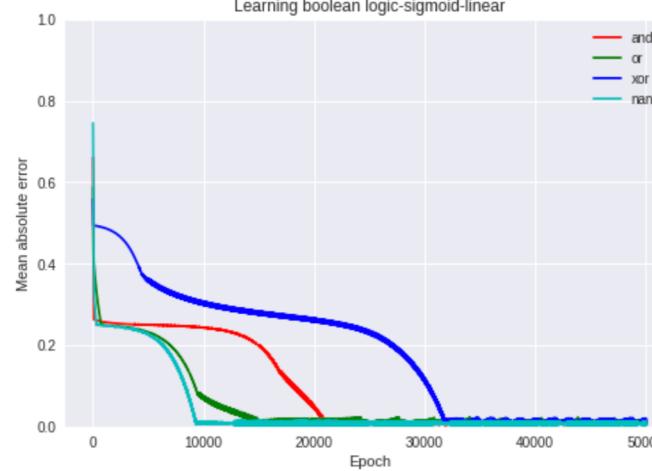
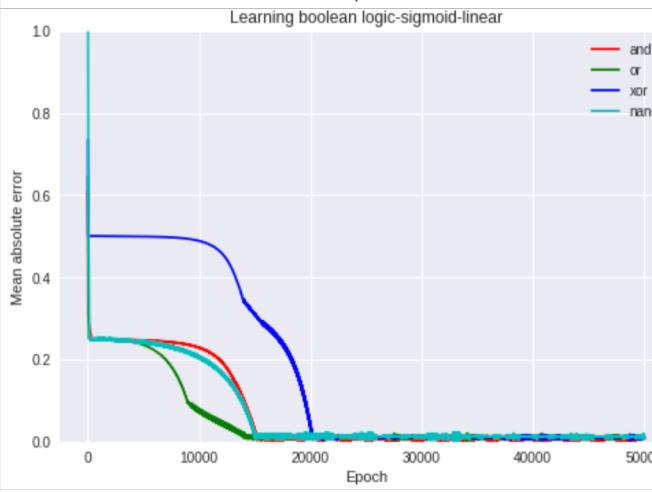
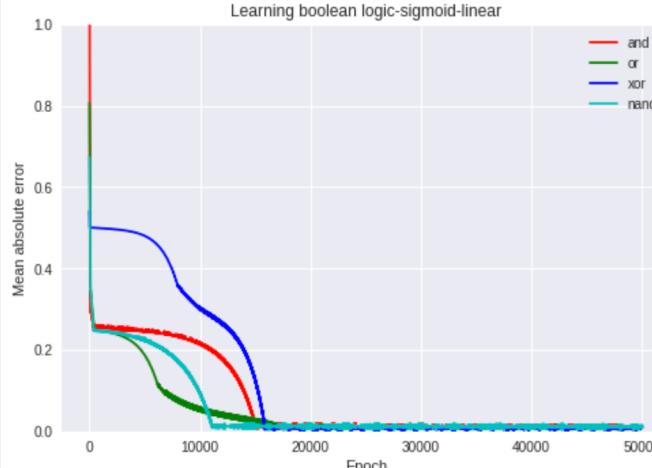
All with:  
sigmoid-linear  
optimizer=adam

# Optimizer matters



All with:  
sigmoid-linear  
batch-size=32

# Random seed matters (at least for sgd)



Three runs, all with:  
sigmoid-linear  
batch-size=32  
Optimizer = sgd