

Learning Systems 2018: Lecture 4 – Tensorflow Overview



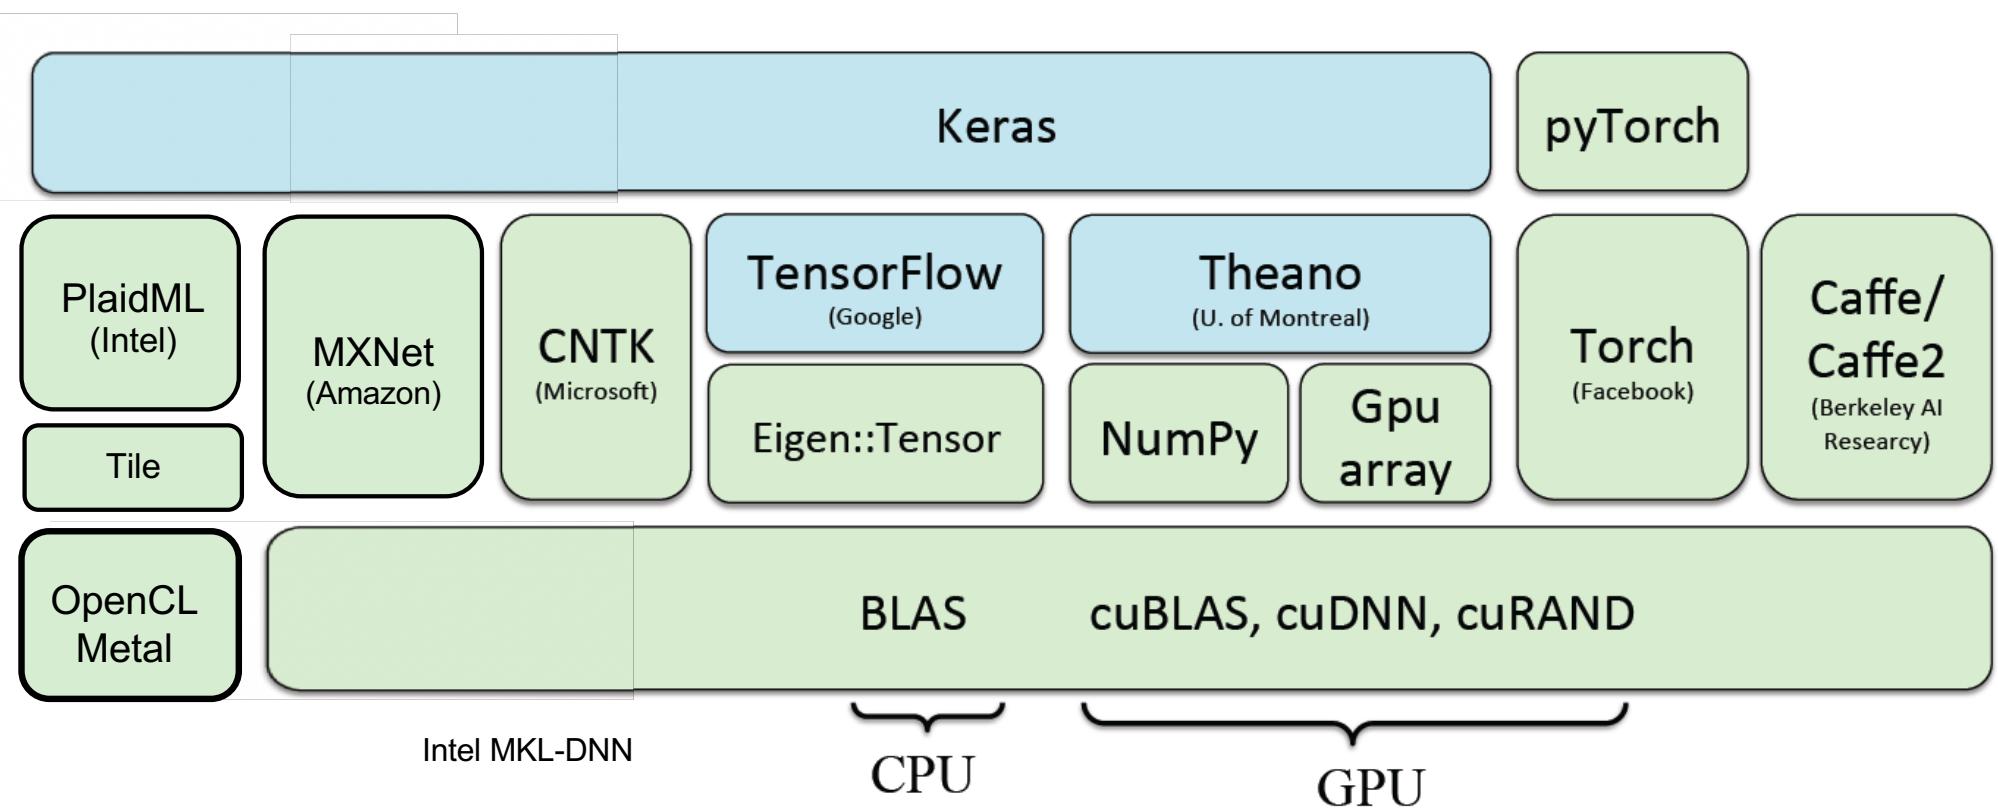
Crescat scientia; vita excolatur

Ian Foster and Rick Stevens
Argonne National Laboratory
The University of Chicago

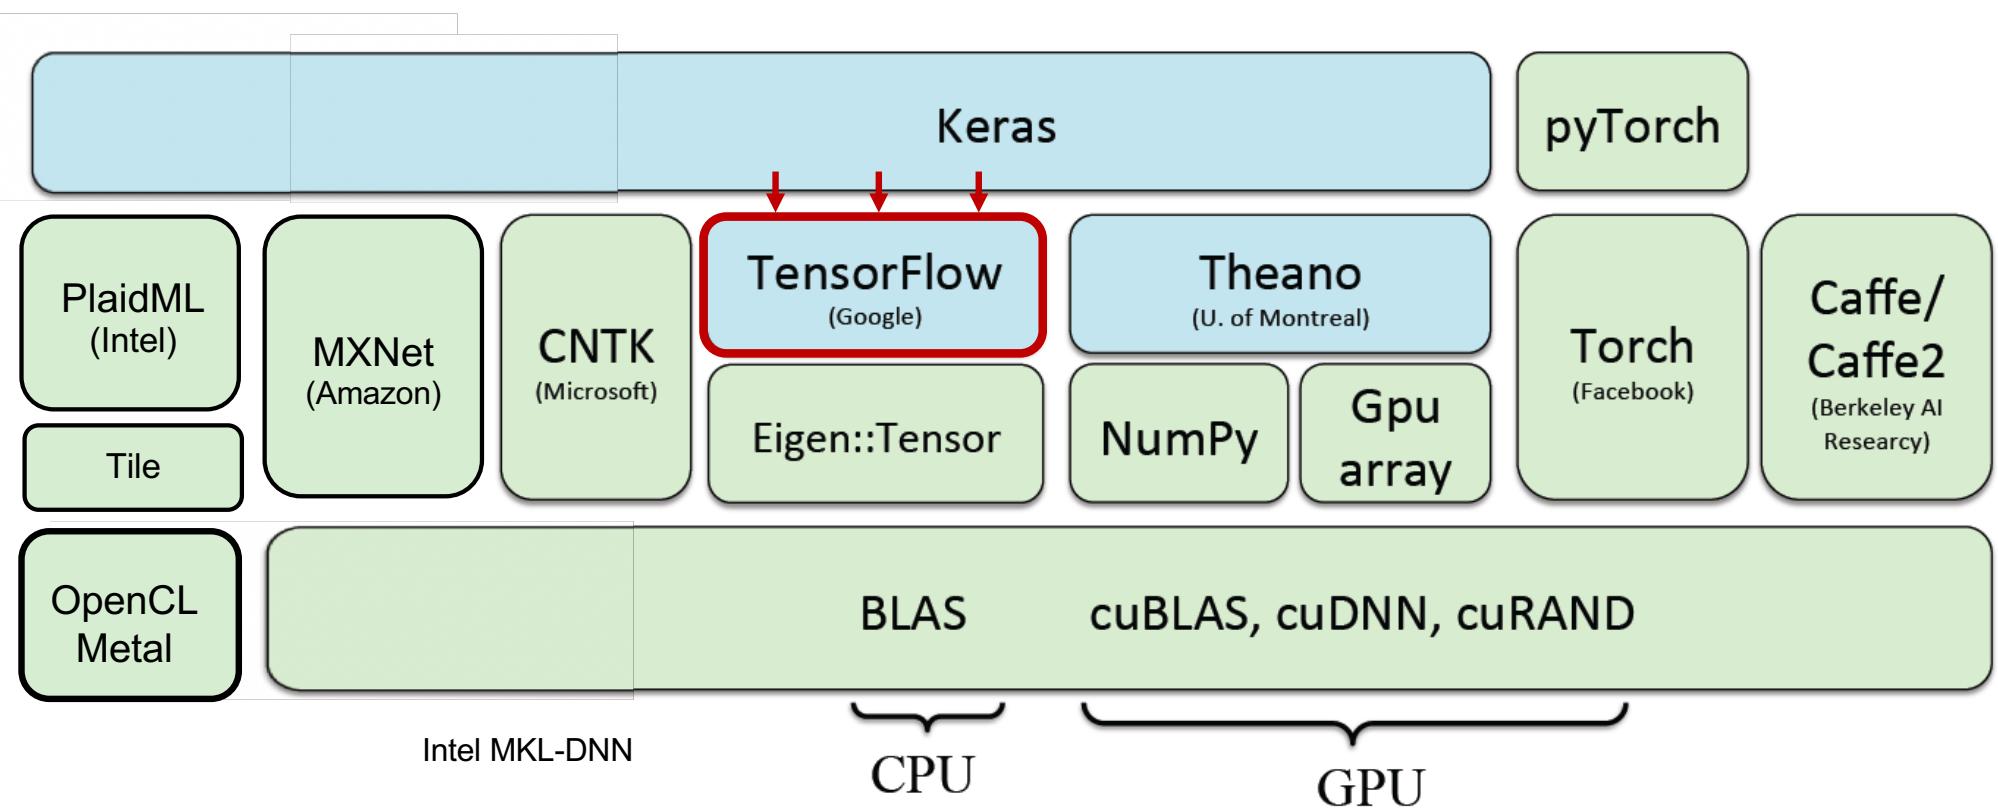
Resources

- <https://medium.com/@the1ju/simple-logistic-regression-using-keras-249e0cc9a970>
- <https://github.com/pbhatnagar3/cs224s-tensorflow-tutorial/blob/master/tensorflow%20MNIST.ipynb>
- <https://web.stanford.edu/class/cs224s/lectures/Tensorflow-tutorial.pdf>
- https://www.tensorflow.org/guide/summaries_and_tensorboard

Deep learning software stacks



Deep learning software stacks

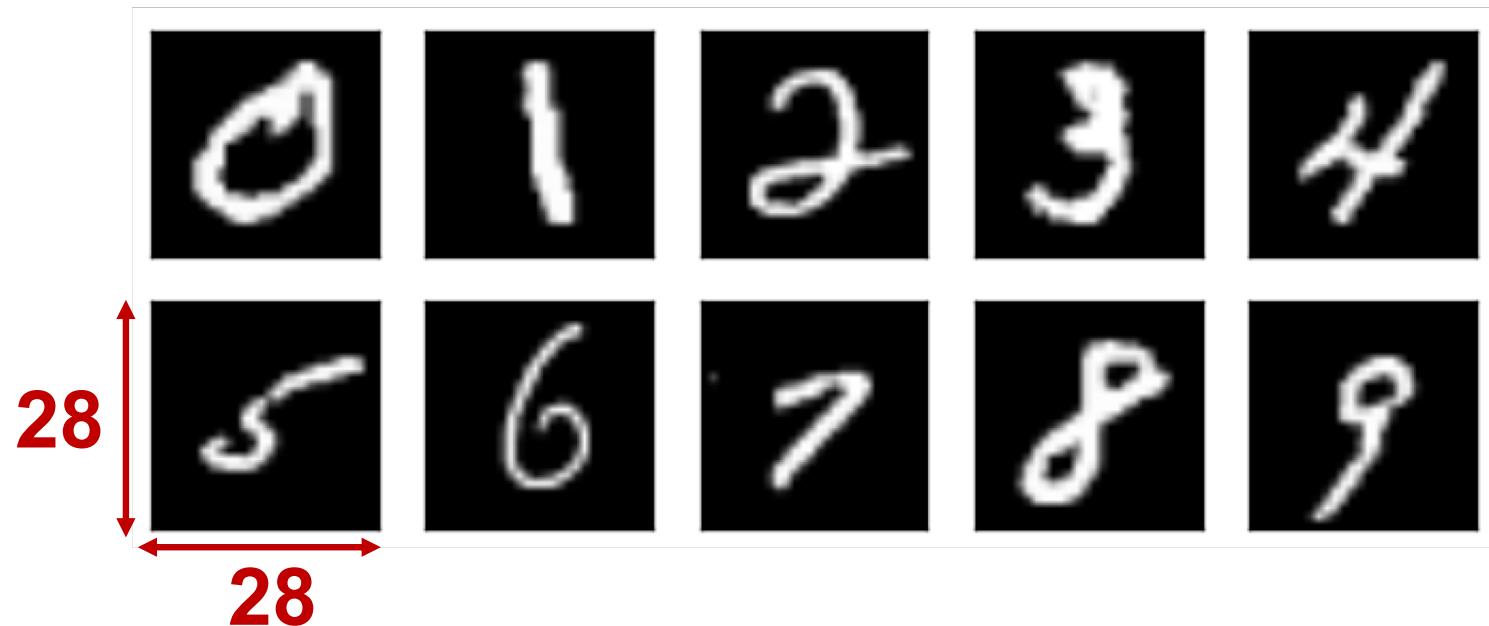


The MNIST test suite

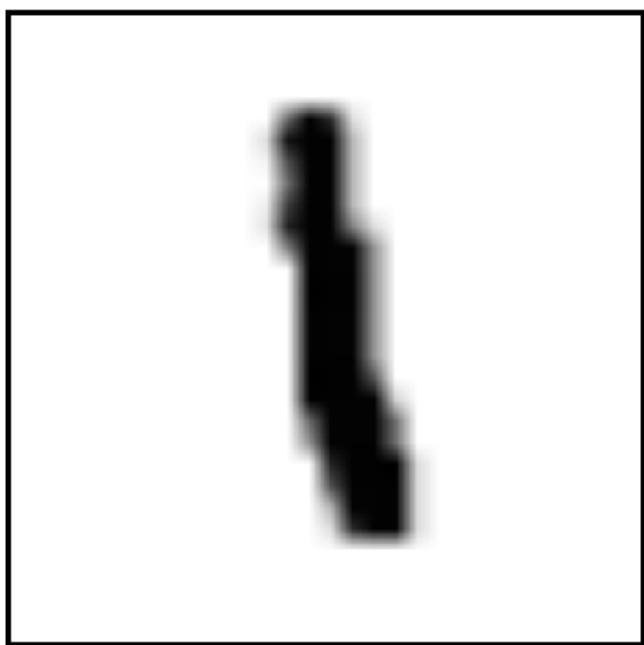


60,000 training examples
10,000 testing examples
Size-normalized
28 x 28 pixels

The MNIST test suite



The MNIST test suite



≈

0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	.6	.8	0	0	0	0	0	0	0	0
0	0	0	0	0	0	.7	1	0	0	0	0	0	0	0	0
0	0	0	0	0	0	.7	1	0	0	0	0	0	0	0	0
0	0	0	0	0	0	.5	1	.4	0	0	0	0	0	0	0
0	0	0	0	0	0	0	1	.4	0	0	0	0	0	0	0
0	0	0	0	0	0	0	1	.4	0	0	0	0	0	0	0
0	0	0	0	0	0	0	1	.7	0	0	0	0	0	0	0
0	0	0	0	0	0	0	1	1	0	0	0	0	0	0	0
0	0	0	0	0	0	0	.9	1	.1	0	0	0	0	0	0
0	0	0	0	0	0	.3	1	.1	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0
0	0	0	0	0	0	0	0	0	0	0	0	0	0	0	0

(showing 14 x 14 for brevity)

Flattened:

[0 0 0 ... 0 .6 .8 0 ... 0 .7 1 0 ... 0 .5 1 .4 0 ... 0 1 .4 0 ... 0 1 .4 0 ... 0 1 .7 0 ... 0 1 1 0 ... 0 .9 1 .1 0 ...]

https://tensorflow.rstudio.com/tensorflow/articles/tutorial_mnist_beginners.html

```
import tensorflow as tf  
mnist = tf.keras.datasets.mnist
```

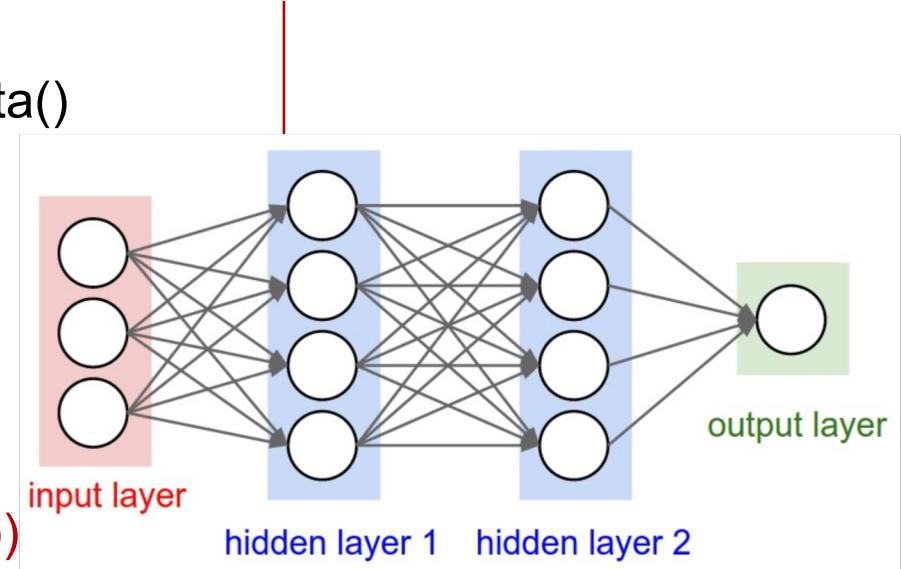
```
(x_train, y_train), (x_test, y_test) = mnist.load_data()  
x_train, x_test = x_train / 255.0, x_test / 255.0
```

```
model = Sequential()  
model.add(Flatten(input_shape=(28, 28)))  
model.add(Dense(units=256, activation='relu'))  
model.add(Dense(units=256, activation='relu'))  
model.add(Dense(units=10, activation='softmax'))
```

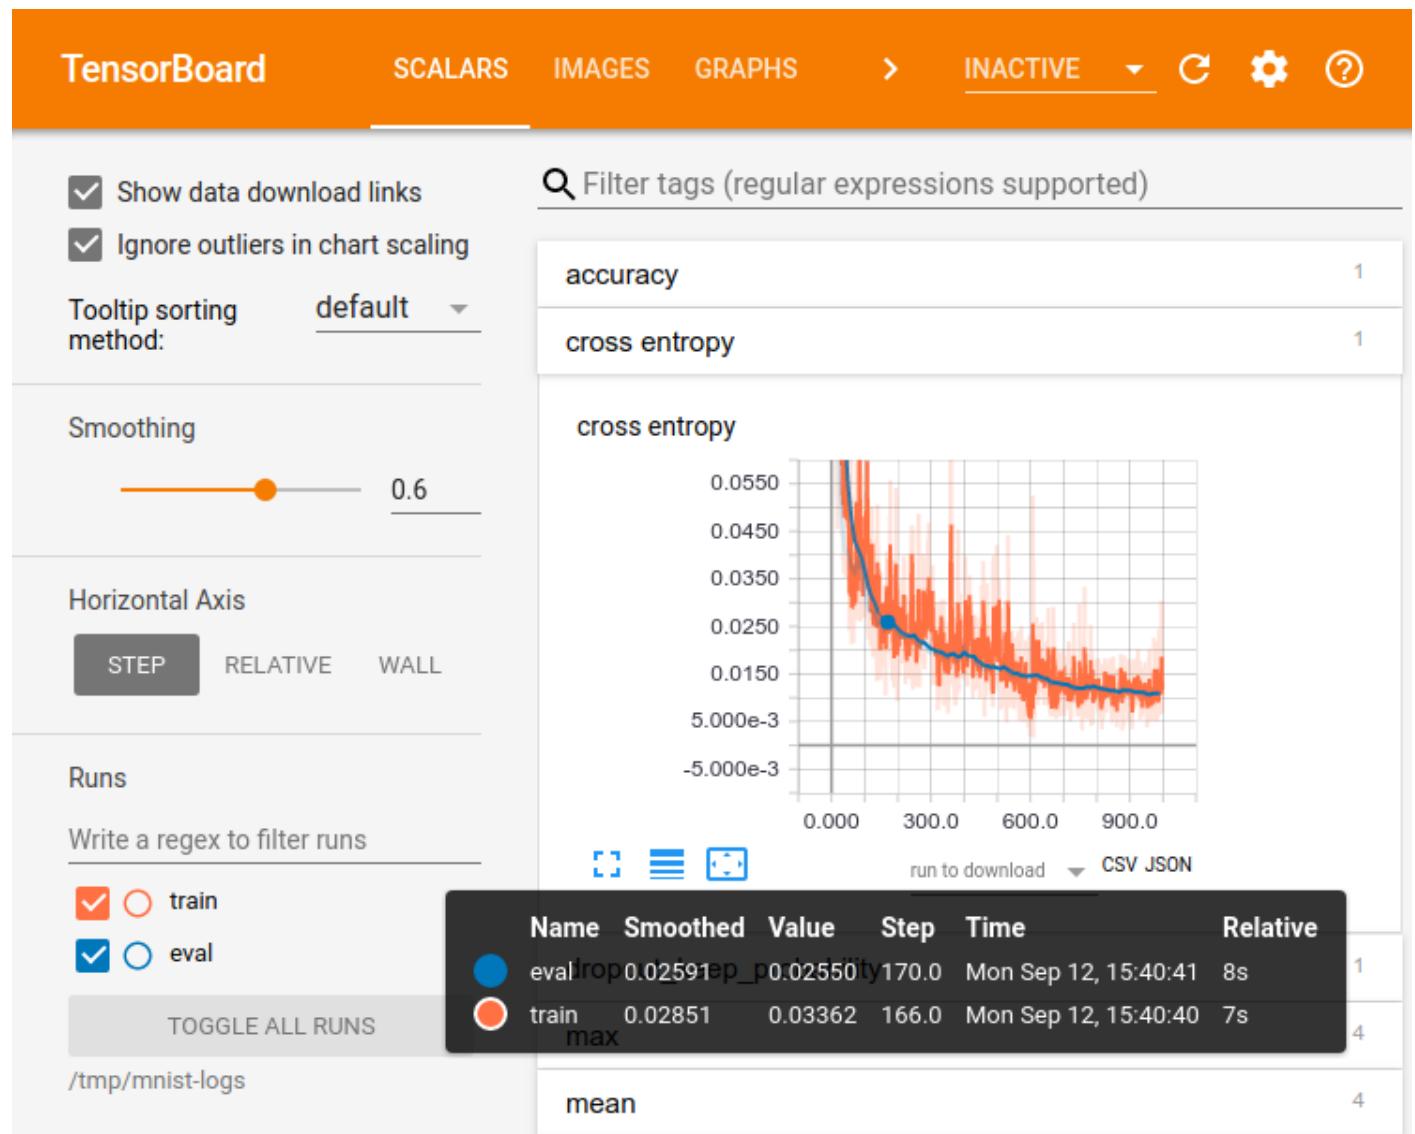
```
model.compile(optimizer='adam',  
              loss='sparse_categorical_crossentropy',  
              metrics=['accuracy'])
```

```
model.fit(x_train, y_train, epochs=5)  
model.evaluate(x_test, y_test)
```

(Simple) MNIST in Keras



Keras and Tensorboard



▼ Set up Tensorboard ([Source](#))

```
[ ] LOG_DIR = '/tmp/log'
get_ipython().system_raw(
    'tensorboard --logdir {} --host 0.0.0.0 --port 6006 &'
    .format(LOG_DIR)
)

[ ] ! wget https://bin.equinox.io/c/4VmDzA7iaHb/ngrok-stable-linux-amd64.zip > /dev/null 2>&1
! yes | unzip ngrok-stable-linux-amd64.zip > /dev/null 2>&1

[ ] get_ipython().system_raw('./ngrok http 6006 &')

[ ] ! curl -s http://localhost:4040/api/tunnels | python3 -c \
    "import sys, json; print(json.load(sys.stdin)['tunnels'][0]['public_url'])"

[ ] ➔ http://8fd226b5.ngrok.io
```

```

# tf Graph input
X = tf.placeholder("float", [None, num_input])
Y = tf.placeholder("float", [None, num_classes])

# Store layers weight & bias
weights = {
    'h1': tf.Variable(tf.random_normal([num_input, n_hidden_1])),
    'h2': tf.Variable(tf.random_normal([n_hidden_1, n_hidden_2])),
    'out': tf.Variable(tf.random_normal([n_hidden_2, num_classes]))
}
biases = {
    'b1': tf.Variable(tf.random_normal([n_hidden_1])),
    'b2': tf.Variable(tf.random_normal([n_hidden_2])),
    'out': tf.Variable(tf.random_normal([num_classes]))
}

# Create model
def neural_net(x):
    # Hidden fully connected layer with 256 neurons
    layer_1 = tf.add(tf.matmul(x, weights['h1']), biases['b1'])
    # Hidden fully connected layer with 256 neurons
    layer_2 = tf.add(tf.matmul(layer_1, weights['h2']), biases['b2'])
    # Output fully connected layer with a neuron for each class
    out_layer = tf.matmul(layer_2, weights['out']) + biases['out']
    return out_layer

# Construct model
logits = neural_net(X)

# Define loss and optimizer
loss_op = tf.reduce_mean(tf.nn.softmax_cross_entropy_with_logits(
    logits=logits, labels=Y))
optimizer = tf.train.AdamOptimizer(learning_rate=learning_rate)
train_op = optimizer.minimize(loss_op)

# Evaluate model (with test logits, for dropout to be disabled)
correct_pred = tf.equal(tf.argmax(logits, 1), tf.argmax(Y, 1))
accuracy = tf.reduce_mean(tf.cast(correct_pred, tf.float32))

# Initialize the variables (i.e. assign their default value)
init = tf.global_variables_initializer()

```

```

# Start training
with tf.Session() as sess:

    # Run the initializer
    sess.run(init)

    for step in range(1, num_steps+1):
        batch_x, batch_y = mnist.train.next_batch(batch_size)
        # Run optimization op (backprop)
        sess.run(train_op, feed_dict={X: batch_x, Y: batch_y})
        if step % display_step == 0 or step == 1:
            # Calculate batch loss and accuracy
            loss, acc = sess.run([loss_op, accuracy], feed_dict={X: batch_x,
                                                                Y: batch_y})
            print("Step " + str(step) + ", Minibatch Loss= " + \
                  "{:.4f}".format(loss) + ", Training Accuracy= " + \
                  "{:.3f}".format(acc))

    print("Optimization Finished!")

    # Calculate accuracy for MNIST test images
    print("Testing Accuracy:", \
          sess.run(accuracy, feed_dict={X: mnist.test.images,
                                       Y: mnist.test.labels}))

```

MNIST in Tensorflow

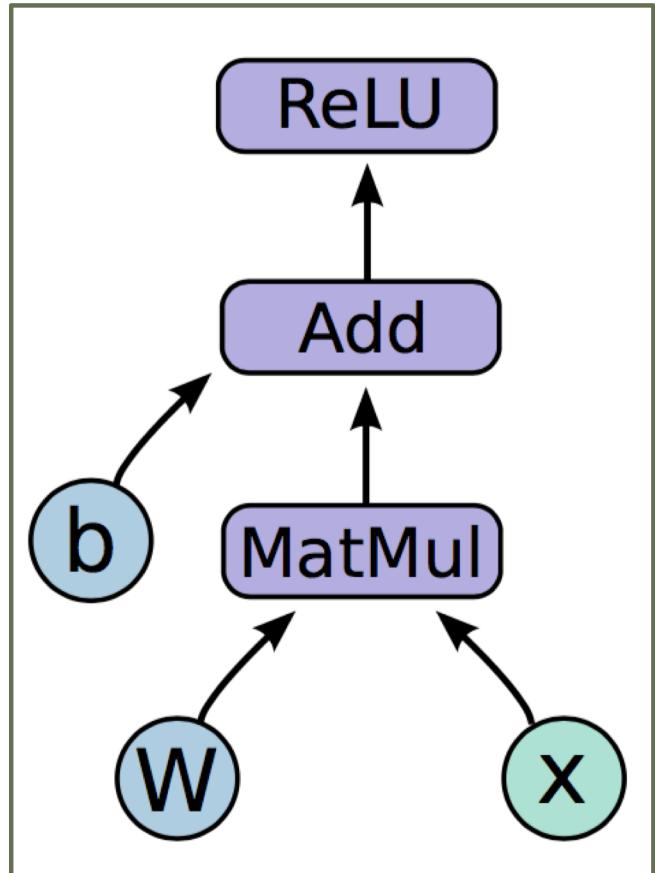


A screenshot of the TensorFlow GitHub repository page. The page shows basic statistics: Watch (8,439), Star (111,471), Fork (68,284). Below these are navigation links: Code (selected), Issues (1,489), Pull requests (257), Projects (0), and Insights. A description of the repository follows: "An Open Source Machine Learning Framework for Everyone" with a link to <https://tensorflow.org>. Below this are several tags: tensorflow, machine-learning, python, deep-learning, deep-neural-networks, neural-network, ml, and distributed. At the bottom, there are metrics: 41,266 commits, 23 branches, 70 releases, 1,665 contributors, and a license Apache-2.0.

- Open source library for numerical computation using data flow graphs
- Originally developed by Google Brain Team for machine learning research
- “Tensorflow is an interface for expressing machine learning algorithms, and an implementation for executing such algorithms” [Abadi et al., 2015]
- “Tensorflow is inordinately complicated” [Internet commentators]

TensorFlow: The basic idea

- 1) We express a computation as a **graph**
 - Graph nodes are **operations** which have any number of inputs and outputs
 - Graph edges are **tensors** [multi-dimensional arrays] which flow between nodes
- 2) Having defined a graph, we can deploy it with a **session**: a binding to a particular execution context (e.g. CPU, GPU)

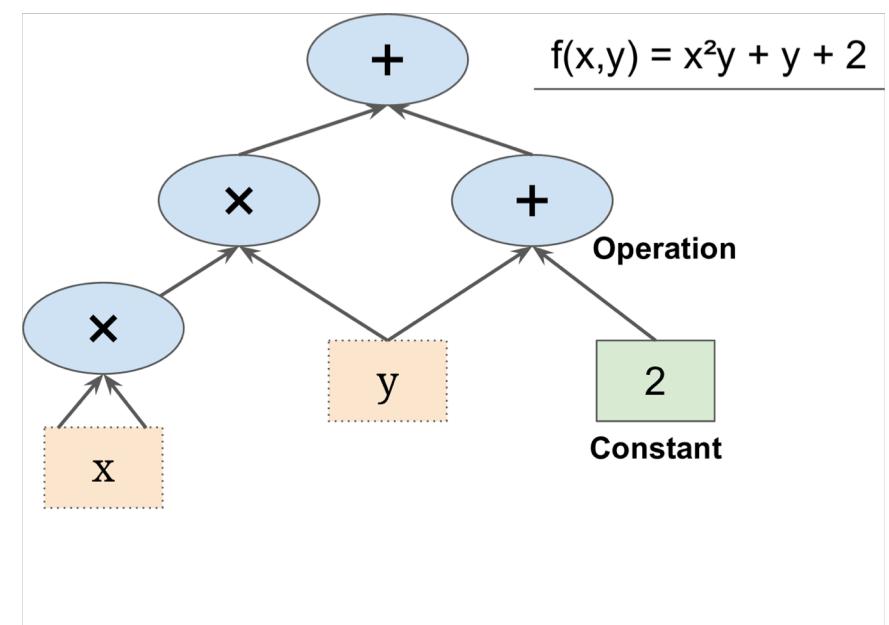


$$h = \text{ReLU}(Wx + b)$$

Tensorflow Graphs: A simple example

Graph:

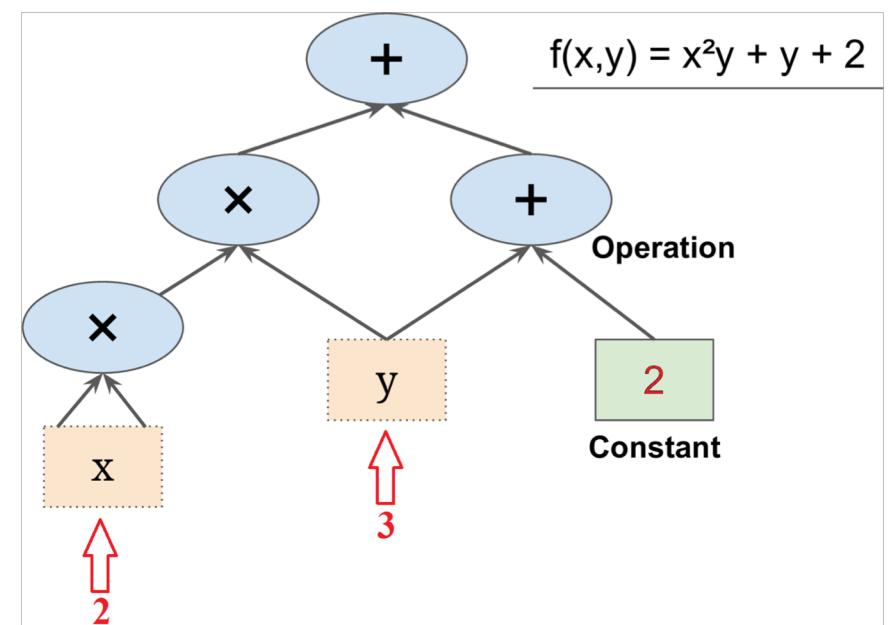
-- Nodes = Operations



Tensorflow Graphs: A simple example

Graph:

- Nodes = Operations
- Edges = Tensors



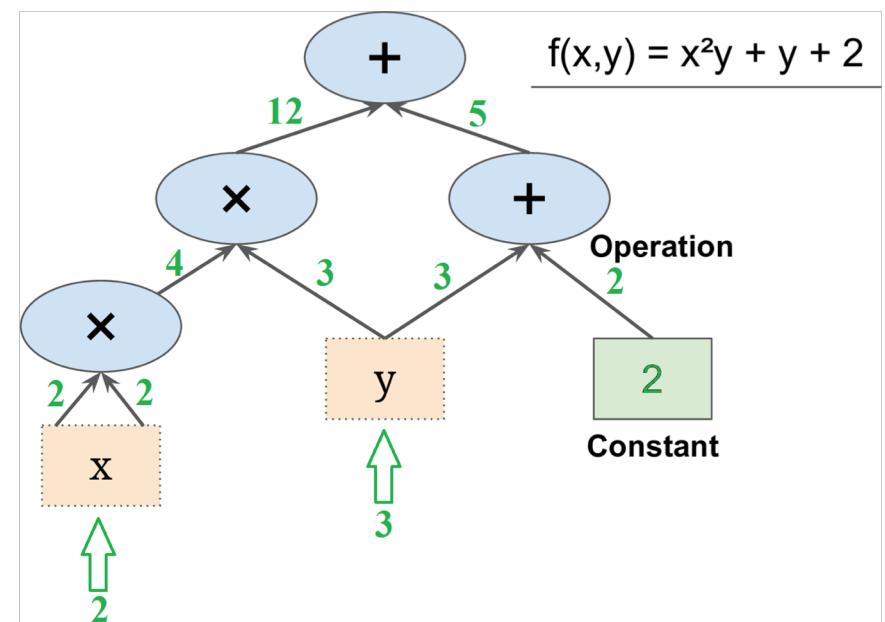
Tensorflow Graphs: A simple example

Graph:

- Nodes = Operations
- Edges = Tensors

Session:

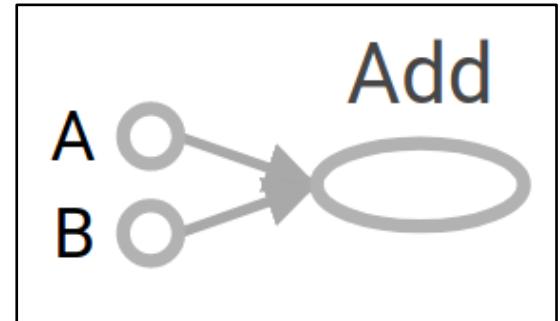
- Tensor = Data
- Tensor + flow = Data + flow



Tensorflow code

```
import tensorflow as tf  
a = tf.constant(2, name='A')  
b = tf.constant(3, name='B')  
c = tf.add(a, b, name='Add')  
  
print(c)
```

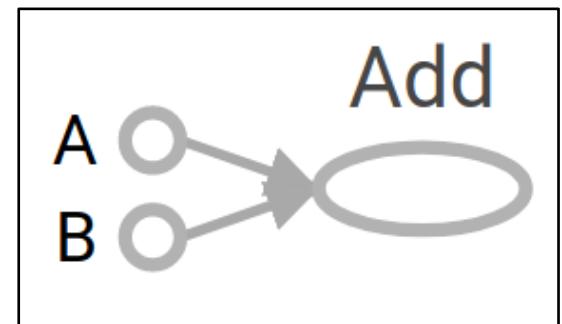
Graph



Tensorflow code

```
import tensorflow as tf  
a = tf.constant(2, name='A')  
b = tf.constant(3, name='B')  
c = tf.add(a, b, name='Add')  
  
print(c)  
  
Tensor("Add:0", shape=(), dtype=int32)
```

Graph



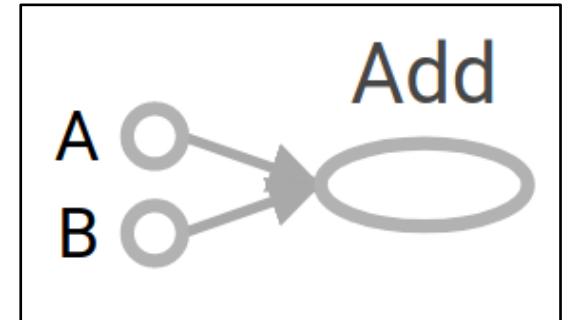
Variables

- a = {Tensor} Tensor("A:0", shape=(), dtype=int32)
- b = {Tensor} Tensor("B:0", shape=(), dtype=int32)
- c = {Tensor} Tensor("Add:0", shape=(), dtype=int32)

Tensorflow code

```
import tensorflow as tf  
a = tf.constant(2, name='A')  
b = tf.constant(3, name='B')  
c = tf.add(a, b, name='Add')  
  
with tf.Session() as sess:  
    print(sess.run(c))  
  
5
```

Graph



Variables

- a = {Tensor} Tensor("A:0", shape=(), dtype=int32)
- b = {Tensor} Tensor("B:0", shape=(), dtype=int32)
- c = {Tensor} Tensor("Add:0", shape=(), dtype=int32)

We:

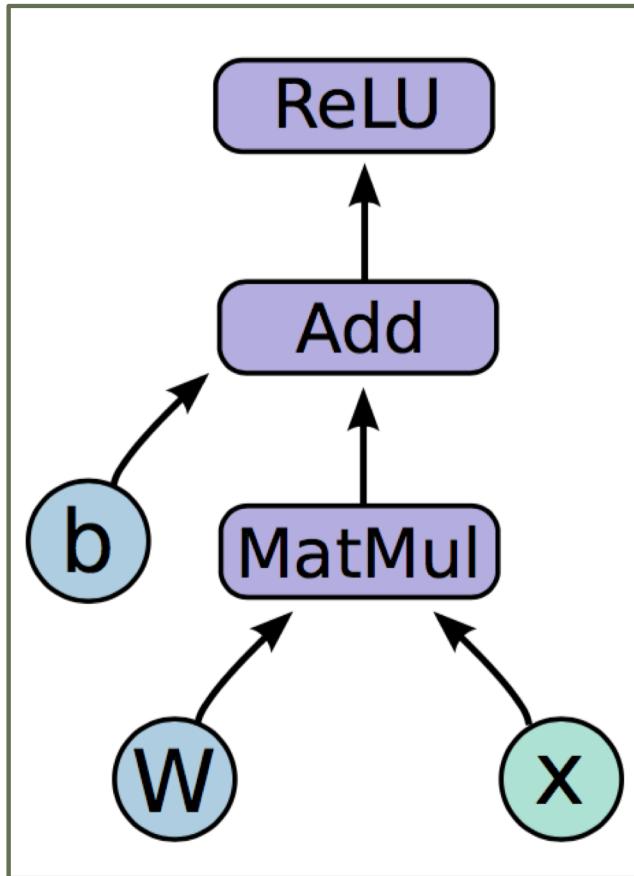
1. Build the GRAPH which represents the data flow of the computation
2. Run the SESSION which executes the operations on the graph

A more complex example

Variables are stateful nodes that output their current value.

State is retained across multiple executions of a graph

Primarily used for parameters



$$h = \text{ReLU}(Wx + b)$$

Mathematical operations:

`MatMul`: Multiply two matrices

`Add`: Add elementwise

`ReLU`: Activate with elementwise rectified linear function

Placeholders are nodes whose value is fed in at execution time

(inputs, labels, ...)

Data types: Constants

Constants create constant values

For example:

```
s = tf.constant(2, name='scalar')
m = tf.constant([[1, 2], [3, 4]], name='matrix')
```

```
tf.constant( value,
            dtype=None,
            shape=None,
            name='Const',
            verify_shape=False
        )
```

Data types: Variables

Variables are stateful nodes that output their current value

- They can be saved and restored
- Gradient updates apply to all variables in the graph

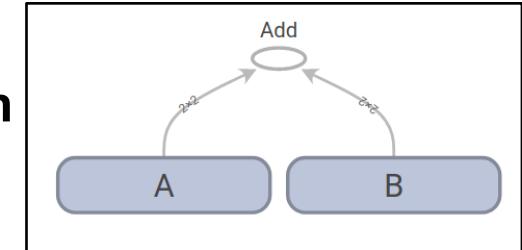
Used for network parameters
(weights and biases)

```
get_variable(  
    name,  
    shape=None,  
    dtype=None,  
    initializer=None,  
    regularizer=None,  
    trainable=True,  
    collections=None,  
    caching_device=None,  
    partitioner=None,  
    validate_shape=True,  
    use_resource=None,  
    custom_getter=None,  
    constraint=None  
)
```

```
s = tf.get_variable(name='scalar2', initializer=tf.constant(2))  
m = tf.get_variable('matrix', initializer=tf.constant([[0, 1], [2, 3]]))  
M = tf.get_variable('big_matrix', shape=(784, 10), initializer=tf.zeros_initializer())  
W = tf.get_variable('weight', shape=(784, 10),  
                    initializer=tf.truncated_normal_initializer(mean=0.0, stddev=0.01))
```

Using variables

Graph



```
# Create graph
a = tf.get_variable(name="A", initializer=tf.constant([[0, 1], [2, 3]]))
b = tf.get_variable(name="B", initializer=tf.constant([[4, 5], [6, 7]]))
c = tf.add(a, b, name="Add")
```

```
# Launch the graph in a session
with tf.Session() as sess:
```

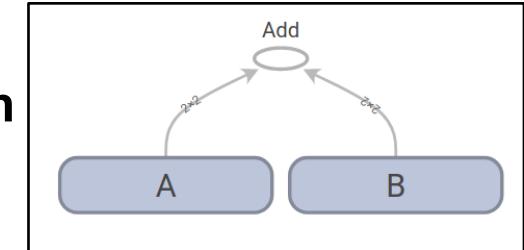
```
# Now we can run the desired operation
print(sess.run(c))
```

Variables

```
■ a = {Variable} <tf.Variable 'A:0' shape=(2, 2) dtype=int32_ref>
■ b = {Variable} <tf.Variable 'B:0' shape=(2, 2) dtype=int32_ref>
■ c = {Tensor} Tensor("Add:0", shape=(2, 2), dtype=int32)
```

Using variables

Graph



```
# Create graph
a = tf.get_variable(name="A", initializer=tf.constant([[0, 1], [2, 3]]))
b = tf.get_variable(name="B", initializer=tf.constant([[4, 5], [6, 7]]))
c = tf.add(a, b, name="Add")
```

```
# Launch the graph in a session
with tf.Session() as sess:
```

Variables

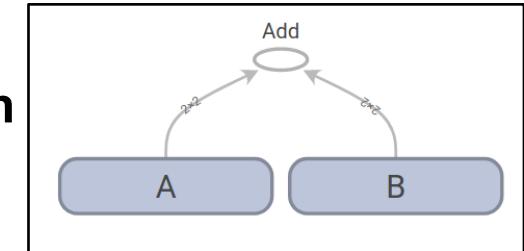
```
■ a = {Variable} <tf.Variable 'A:0' shape=(2, 2) dtype=int32_ref>
■ b = {Variable} <tf.Variable 'B:0' shape=(2, 2) dtype=int32_ref>
■ c = {Tensor} Tensor("Add:0", shape=(2, 2), dtype=int32)
```

```
# Now we can run the desired operation
print(sess.run(c))
```

FailedPreconditionError: Attempting to use uninitialized value

Using variables

Graph



```
# Create graph
a = tf.get_variable(name="A", initializer=tf.constant([[0, 1], [2, 3]]))
b = tf.get_variable(name="B", initializer=tf.constant([[4, 5], [6, 7]]))
c = tf.add(a, b, name="Add")
# Add an Op to initialize variables
init_op = tf.global_variables_initializer()
# Launch the graph in a session
with tf.Session() as sess:
    # Run the variable initializer
    sess.run(init_op)
    # Now we can run the desired operation
    print(sess.run(c))
[[ 4,  6], [ 8, 10]]
```

Variables

```
■ a = {Variable} <tf.Variable 'A:0' shape=(2, 2) dtype=int32_ref>
■ b = {Variable} <tf.Variable 'B:0' shape=(2, 2) dtype=int32_ref>
■ c = {Tensor} Tensor("Add:0", shape=(2, 2), dtype=int32)
```

Data types: Placeholder

```
tf.placeholder( dtype,  
                shape=None,  
                name=None  
            )
```

A **Placeholder** is a node whose value is fed in at execution time.

- Assemble the graph without knowing the values needed for computation
- Supply the data at execution time.

Used for input data (in classification task: Inputs and labels)

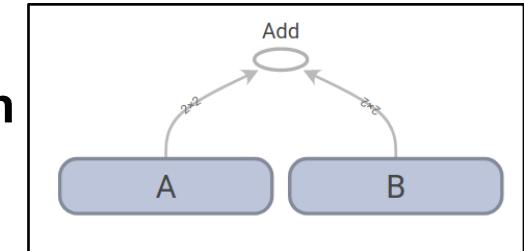
```
a = tf.placeholder(tf.float32, shape=[5])  
b = tf.placeholder(dtype=tf.float32, shape=None, name=None)  
X = tf.placeholder(tf.float32, shape=[None, 784], name='input')  
Y = tf.placeholder(tf.float32, shape=[None, 10], name='label')
```

Using placeholders

```
import tensorflow as tf
a = tf.constant([5, 5, 5], tf.float32, name='A')
b = tf.placeholder(tf.float32, shape=[3], name='B')
c = tf.add(a, b, name="Add")

with tf.Session() as sess:
    print(sess.run(c))
```

Graph



Variables

```
☰ a = {Tensor} Tensor("A:0", shape=(3,), dtype=float32)
☰ b = {Tensor} Tensor("B:0", shape=(3,), dtype=float32)
☰ c = {Tensor} Tensor("Add:0", shape=(3,), dtype=float32)
```

Using placeholders

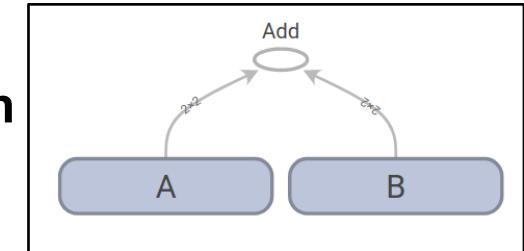
```
import tensorflow as tf  
a = tf.constant([5, 5, 5], tf.float32, name='A')  
b = tf.placeholder(tf.float32, shape=[3], name='B')  
c = tf.add(a, b, name="Add")
```

with tf.Session() as sess:

```
print(sess.run(c))
```

You must feed a value for placeholder tensor
'B' with dtype float and shape [3]

Graph



Variables

```
■ a = {Tensor} Tensor("A:0", shape=(3,), dtype=float32)  
■ b = {Tensor} Tensor("B:0", shape=(3,), dtype=float32)  
■ c = {Tensor} Tensor("Add:0", shape=(3,), dtype=float32)
```

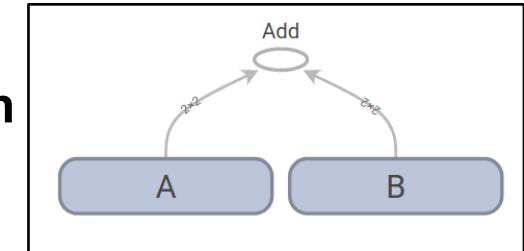
Using placeholders

```
import tensorflow as tf
a = tf.constant([5, 5, 5], tf.float32, name='A')
b = tf.placeholder(tf.float32, shape=[3], name='B')
c = tf.add(a, b, name="Add")
```

```
with tf.Session() as sess:
    # Create a dictionary:
    d = {b: [1, 2, 3]}
    # Feed it to the placeholder
    print(sess.run(c, feed_dict=d))
```

[6, 7, 8]

Graph

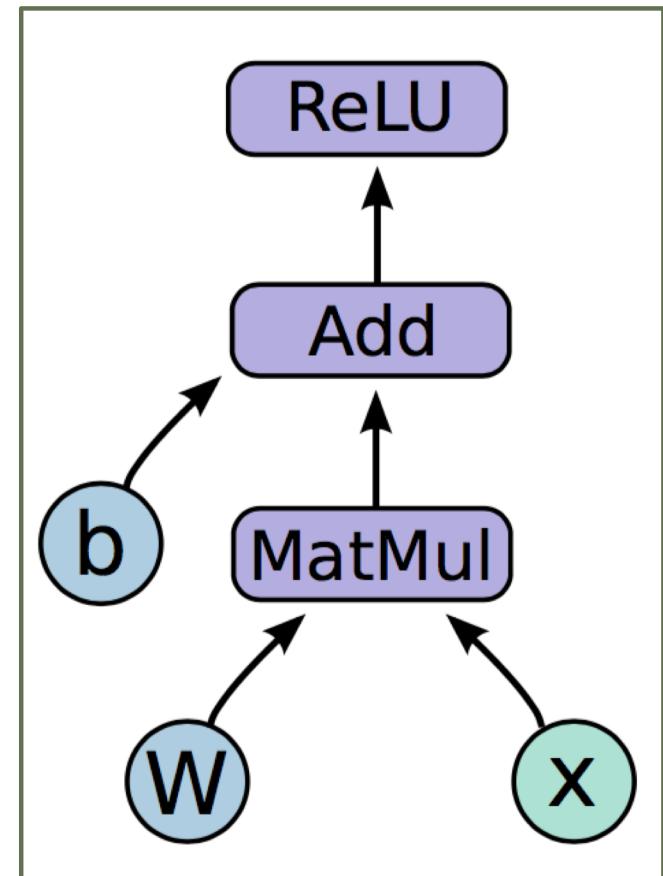


Variables

```
■ a = {Tensor} Tensor("A:0", shape=(3,), dtype=float32)
■ b = {Tensor} Tensor("B:0", shape=(3,), dtype=float32)
■ c = {Tensor} Tensor("Add:0", shape=(3,), dtype=float32)
■ d = {dict} {<tf.Tensor 'B:0' shape=(3,) dtype=float32>: [1, 2, 3]}
```

Back to our example: Creating a graph

```
import tensorflow as tf  
  
b = tf.Variable(tf.zeros((100,)))  
W = tf.Variable(tf.random_uniform((784, 100), -1, 1))  
  
x = tf.placeholder(tf.float32, (100, 784))  
  
h = tf.nn.relu(tf.matmul(x, W) + b)
```



$$h = \text{ReLU}(Wx + b)$$

Creating a graph

```
import tensorflow as tf

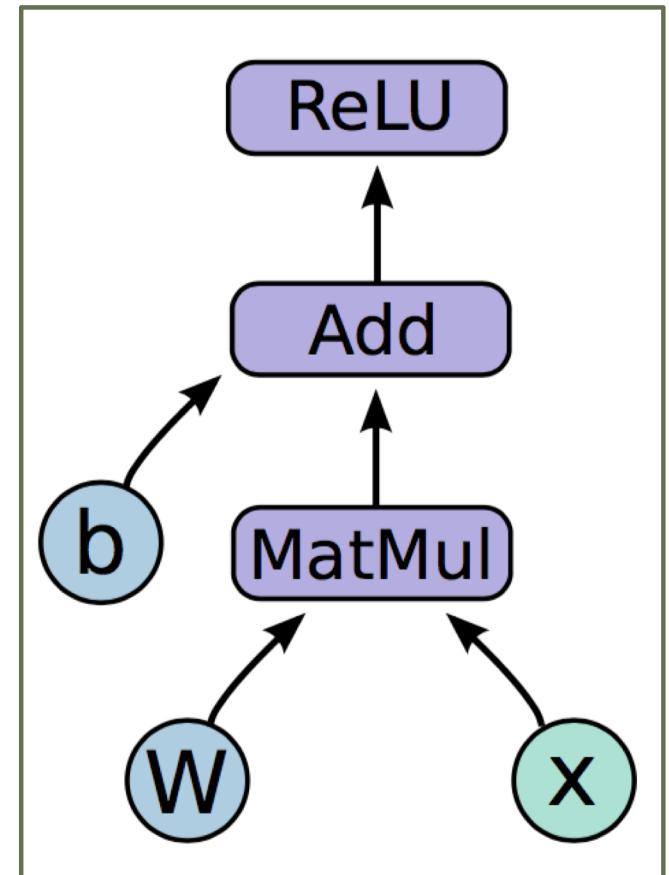
b = tf.Variable(tf.zeros((100,)))
W = tf.Variable(tf.random_uniform((784, 100), -1, 1))

x = tf.placeholder(tf.float32, (100, 784))

h = tf.nn.relu(tf.matmul(x, W) + b)
```

Deploying the graph into a session

```
with tf.Session() as sess:
    sess.run(tf.initialize_all_variables())
    sess.run(h, {x: np.random.random((100, 784))})
```



$$h = \text{ReLU}(Wx + b)$$

```

# tf Graph input
X = tf.placeholder("float", [None, num_input])
Y = tf.placeholder("float", [None, num_classes])

# Store layers weight & bias
weights = {
    'h1': tf.Variable(tf.random_normal([num_input, n_hidden_1])),
    'h2': tf.Variable(tf.random_normal([n_hidden_1, n_hidden_2])),
    'out': tf.Variable(tf.random_normal([n_hidden_2, num_classes]))
}
biases = {
    'b1': tf.Variable(tf.random_normal([n_hidden_1])),
    'b2': tf.Variable(tf.random_normal([n_hidden_2])),
    'out': tf.Variable(tf.random_normal([num_classes]))
}

# Create model
def neural_net(x):
    # Hidden fully connected layer with 256 neurons
    layer_1 = tf.add(tf.matmul(x, weights['h1']), biases['b1'])
    # Hidden fully connected layer with 256 neurons
    layer_2 = tf.add(tf.matmul(layer_1, weights['h2']), biases['b2'])
    # Output fully connected layer with a neuron for each class
    out_layer = tf.matmul(layer_2, weights['out']) + biases['out']
    return out_layer

# Construct model
logits = neural_net(X)

# Define loss and optimizer
loss_op = tf.reduce_mean(tf.nn.softmax_cross_entropy_with_logits=logits, labels=Y))
optimizer = tf.train.AdamOptimizer(learning_rate=learn_rate)
train_op = optimizer.minimize(loss_op)

# Evaluate model (with test logits, for dropout to be
correct_pred = tf.equal(tf.argmax(logits, 1), tf.argmax(Y, 1))
accuracy = tf.reduce_mean(tf.cast(correct_pred, tf.float32))

# Initialize the variables (i.e. assign their default value)
init = tf.global_variables_initializer()

```

```

# Start training
with tf.Session() as sess:
    # Run the initializer
    sess.run(init)

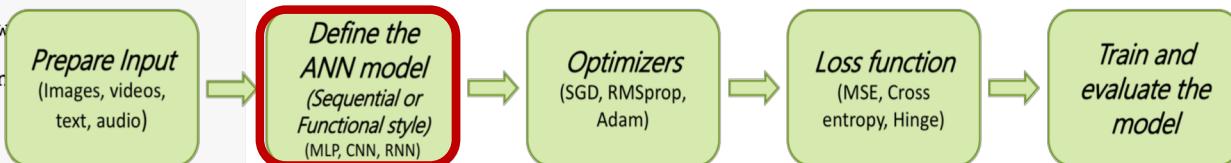
    for step in range(1, num_steps+1):
        batch_x, batch_y = mnist.train.next_batch(batch_size)
        # Run optimization op (backprop)
        sess.run(train_op, feed_dict={X: batch_x, Y: batch_y})
        if step % display_step == 0 or step == 1:
            # Calculate batch loss and accuracy
            loss, acc = sess.run([loss_op, accuracy], feed_dict={X: batch_x,
                                                                Y: batch_y})
            print("Step " + str(step) + ", Minibatch Loss= " + \
                  "{:.4f}".format(loss) + ", Training Accuracy= " + \
                  "{:.3f}".format(acc))

    print("Optimization Finished!")

    # Calculate accuracy for MNIST test images
    print("Testing Accuracy:", \
          sess.run(accuracy, feed_dict={X: mnist.test.images,
                                       Y: mnist.test.labels}))

```

MNIST in Tensorflow



```
# tf Graph input
X = tf.placeholder("float", [None, num_input])
Y = tf.placeholder("float", [None, num_classes])

# Store layers weight & bias
weights = {
    'h1': tf.Variable(tf.random_normal([num_input, n_hidden_1])),
    'h2': tf.Variable(tf.random_normal([n_hidden_1, n_hidden_2])),
    'out': tf.Variable(tf.random_normal([n_hidden_2, num_classes])))
}
biases = {
    'b1': tf.Variable(tf.random_normal([n_hidden_1])),
    'b2': tf.Variable(tf.random_normal([n_hidden_2])),
    'out': tf.Variable(tf.random_normal([num_classes])))
}

# Create model
def neural_net(x):
    # Hidden fully connected layer with 256 neurons
    layer_1 = tf.add(tf.matmul(x, weights['h1']), biases['b1'])
    # Hidden fully connected layer with 256 neurons
    layer_2 = tf.add(tf.matmul(layer_1, weights['h2']), biases['b2'])
    # Output fully connected layer with a neuron for each class
    out_layer = tf.matmul(layer_2, weights['out']) + biases['out']
    return out_layer

# Construct model
logits = neural_net(X)
```

Next up, we add a loss node

Use a **placeholder** for **labels**, as they will be supplied later

Build loss node using labels and **prediction**

$$f_i(\vec{x}) = \frac{e^{x_i}}{\sum_{j=1}^J e^{x_j}} \quad \text{for } i = 1, \dots, J$$

```
prediction = tf.nn.softmax(...) # Output of neural network
```

```
label = tf.placeholder(tf.float32, [100, 10])
```

```
cross_entropy = -tf.reduce_sum(label * tf.log(prediction), axis=1)
```

SoftMax:
Converts scores
to probabilities
that sum to 1

“In information theory, the **cross entropy** between two probability distributions **p** and **q** over the same underlying set of events measures the average number of bits needed to identify an event drawn from the set, if a coding scheme is used that is optimized for an “unnatural” probability distribution **q**, rather than the “true” distribution **p**

$$H(p, q) = - \sum_x p(x) \log q(x)$$

We next want to compute gradients

```
train_step = tf.train.GradientDescentOptimizer(0.5).minimize(cross_entropy)
```

tf.train.GradientDescentOptimizer is an **Optimizer** object

`tf.train.GradientDescentOptimizer(lr).minimize(cross_entropy)` adds optimization **operation** to computation graph

TensorFlow graph **nodes** have **attached gradient operations**

Gradient with respect to **parameters** computed with **backpropagation**

Create the train_step op

```
prediction = tf.nn.softmax(...)  
label = tf.placeholder(tf.float32, [None, 10])  
  
cross_entropy = tf.reduce_mean(-tf.reduce_sum(label * tf.log(prediction),  
                                         reduction_indices=[1]))  
  
train_step = tf.train.GradientDescentOptimizer(0.5).minimize(cross_entropy)
```

Adding the **train** line adds an optimization node to the graph.

Evaluating this “train_step” Python variable at runtime will automatically apply gradients to all of the variables in our graph.

Now we can train the model

We can now train the `train_step` node. To do this we:

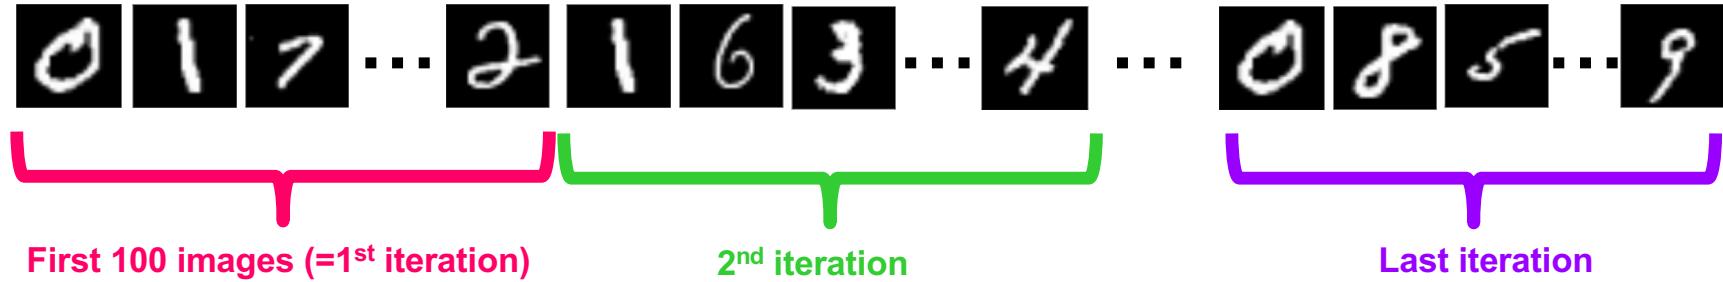
1. Create Session
2. Build training schedule (iterative, feeds in data and labels, trains model)
3. Run `train_step`

```
sess = tf.Session()  
sess.run(tf.initialize_all_variables())  
  
for i in range(1000):  
    batch_x, batch_label = data.next_batch()  
    sess.run(train_step, feed_dict={x: batch_x, label: batch_label})
```

Batches, epochs, and iteration

Example: MNIST data

- Number of training data: **N=60,000**
- Let's take batch size of **B= 100**



How many iterations per epoch? $60,000/100 = 600$ **1 epoch = 600 iterations**

```

# tf Graph input
X = tf.placeholder("float", [None, num_input])
Y = tf.placeholder("float", [None, num_classes])

Input X and Y: placeholders

# Store layers weight & bias
weights = {
    'h1': tf.Variable(tf.random_normal([num_input, n_hidden_1])),
    'h2': tf.Variable(tf.random_normal([n_hidden_1, n_hidden_2])),
    'out': tf.Variable(tf.random_normal([n_hidden_2, num_classes]))
}
biases = {
    'b1': tf.Variable(tf.random_normal([n_hidden_1])),
    'b2': tf.Variable(tf.random_normal([n_hidden_2])),
    'out': tf.Variable(tf.random_normal([num_classes]))
}

```

Weights and biases: Variables

```

# Create model
def neural_net(x):
    # Hidden fully connected layer with 256 neurons
    layer_1 = tf.add(tf.matmul(x, weights['h1']), biases['b1'])
    # Hidden fully connected layer with 256 neurons
    layer_2 = tf.add(tf.matmul(layer_1, weights['h2']), biases['b2'])
    # Output fully connected layer with a neuron for each class
    out_layer = tf.matmul(layer_2, weights['out']) + biases['out']
    return out_layer

```

Model: Three layers

```

# Construct model
logits = neural_net(X)

# Define loss and optimizer
loss_op = tf.reduce_mean(tf.nn.softmax_cross_entropy_with_logits(
    logits=logits, labels=Y))
optimizer = tf.train.AdamOptimizer(learning_rate=learning_rate)
train_op = optimizer.minimize(loss_op)

# Evaluate model (with test logits, for dropout to be disabled)
correct_pred = tf.equal(tf.argmax(logits, 1), tf.argmax(Y, 1))
accuracy = tf.reduce_mean(tf.cast(correct_pred, tf.float32))

# Initialize the variables (i.e. assign their default value)
init = tf.global_variables_initializer()

```

Define loss and training op

```

# Start training
with tf.Session() as sess:
    # Run the initializer
    sess.run(init)

    for step in range(1, num_steps+1):
        batch_x, batch_y = mnist.train.next_batch(batch_size)
        # Run optimization op (backprop)
        sess.run(train_op, feed_dict={X: batch_x, Y: batch_y})
        if step % display_step == 0 or step == 1:
            # Calculate batch loss and accuracy
            loss, acc = sess.run([loss_op, accuracy], feed_dict={X: batch_x,
                                                                Y: batch_y})
            print("Step " + str(step) + ", Minibatch Loss= " + \
                  "{:.4f}".format(loss) + ", Training Accuracy= " + \
                  "{:.3f}".format(acc))

    print("Optimization Finished!")

# Calculate accuracy for MNIST test images
print("Testing Accuracy:", \
      sess.run(accuracy, feed_dict={X: mnist.test.images,
                                   Y: mnist.test.labels}))

```

Train model

Test accuracy

MNIST in Tensorflow



Variable sharing

```
variables_dict = {
    "weights": tf.Variable(tf.random_normal([782, 100]), name="weights")
    "biases": tf.Variable(tf.zeros([100]), name="biases")
}
```

tf.variable_scope() provides simple name-spacing to avoid clashes
tf.get_variable() creates/accesses variables from within a variable scope

```
with tf.variable_scope("foo"):
    v = tf.get_variable("v", shape=[1]) # v.name == "foo/v:0"
```

```
with tf.variable_scope("foo", reuse=True):
    v1 = tf.get_variable("v") # Shared variable found!
```

```
with tf.variable_scope("foo", reuse=False):
    v1 = tf.get_variable("v") # CRASH foo/v:0 already exists!
```

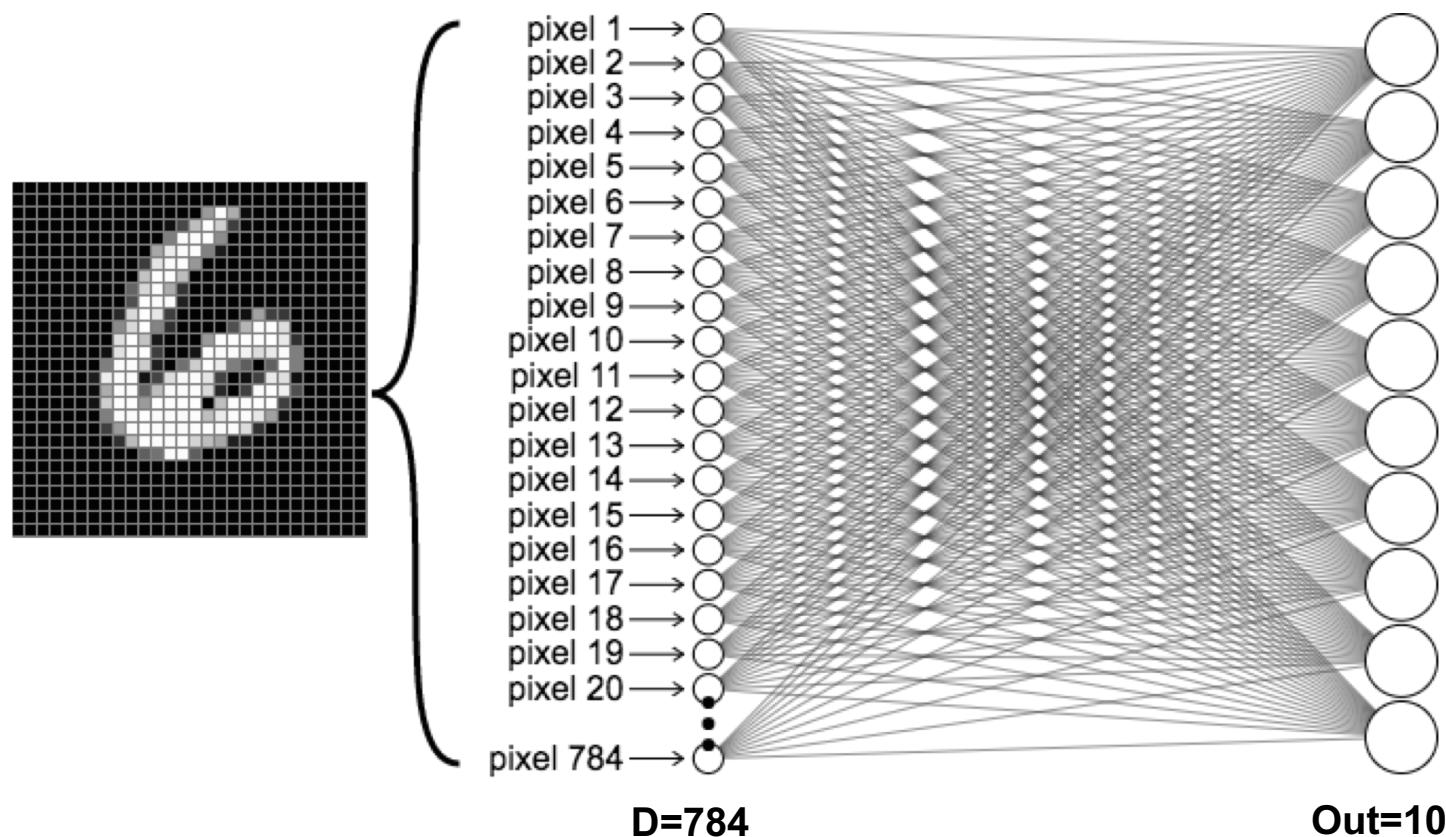
```
def nn_layer(input_tensor, input_dim, output_dim, layer_name, act=tf.nn.relu):
    """Reusable code for making a simple neural net layer.

    It does a matrix multiply, bias add, and then uses ReLU to nonlinearize.
    It also sets up name scoping so that the resultant graph is easy to read,
    and adds a number of summary ops.
    """
    # Adding a name scope ensures logical grouping of the layers in the graph.
    with tf.name_scope(layer_name):
        # This Variable will hold the state of the weights for the layer
        with tf.name_scope('weights'):
            weights = weight_variable([input_dim, output_dim])
            variable_summaries(weights)
        with tf.name_scope('biases'):
            biases = bias_variable([output_dim])
            variable_summaries(biases)
        with tf.name_scope('Wx_plus_b'):
            preactivate = tf.matmul(input_tensor, weights) + biases
            tf.summary.histogram('pre_activations', preactivate)
        activations = act(preactivate, name='activation')
        tf.summary.histogram('activations', activations)
    return activations

hidden1 = nn_layer(x, 784, 500, 'layer1')
```

Back to MNIST

Logistic (i.e., linear) classifier



Logistic classifier

```
input_dim = 28*28
output_dim = nb_classes = 10

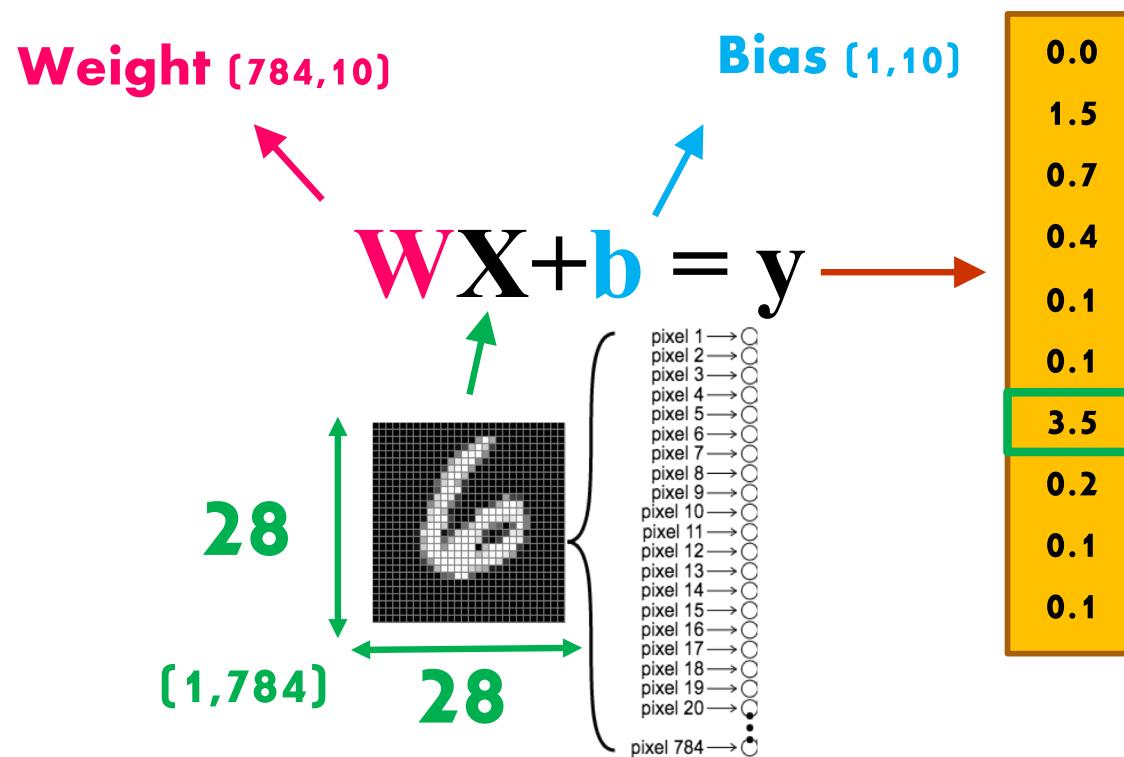
model = Sequential()
model.add(Flatten(input_shape=(28, 28)))
model.add(Dense(output_dim, activation='softmax'))

model.compile(optimizer='adam',
              loss='sparse_categorical_crossentropy',
              metrics=['accuracy'])

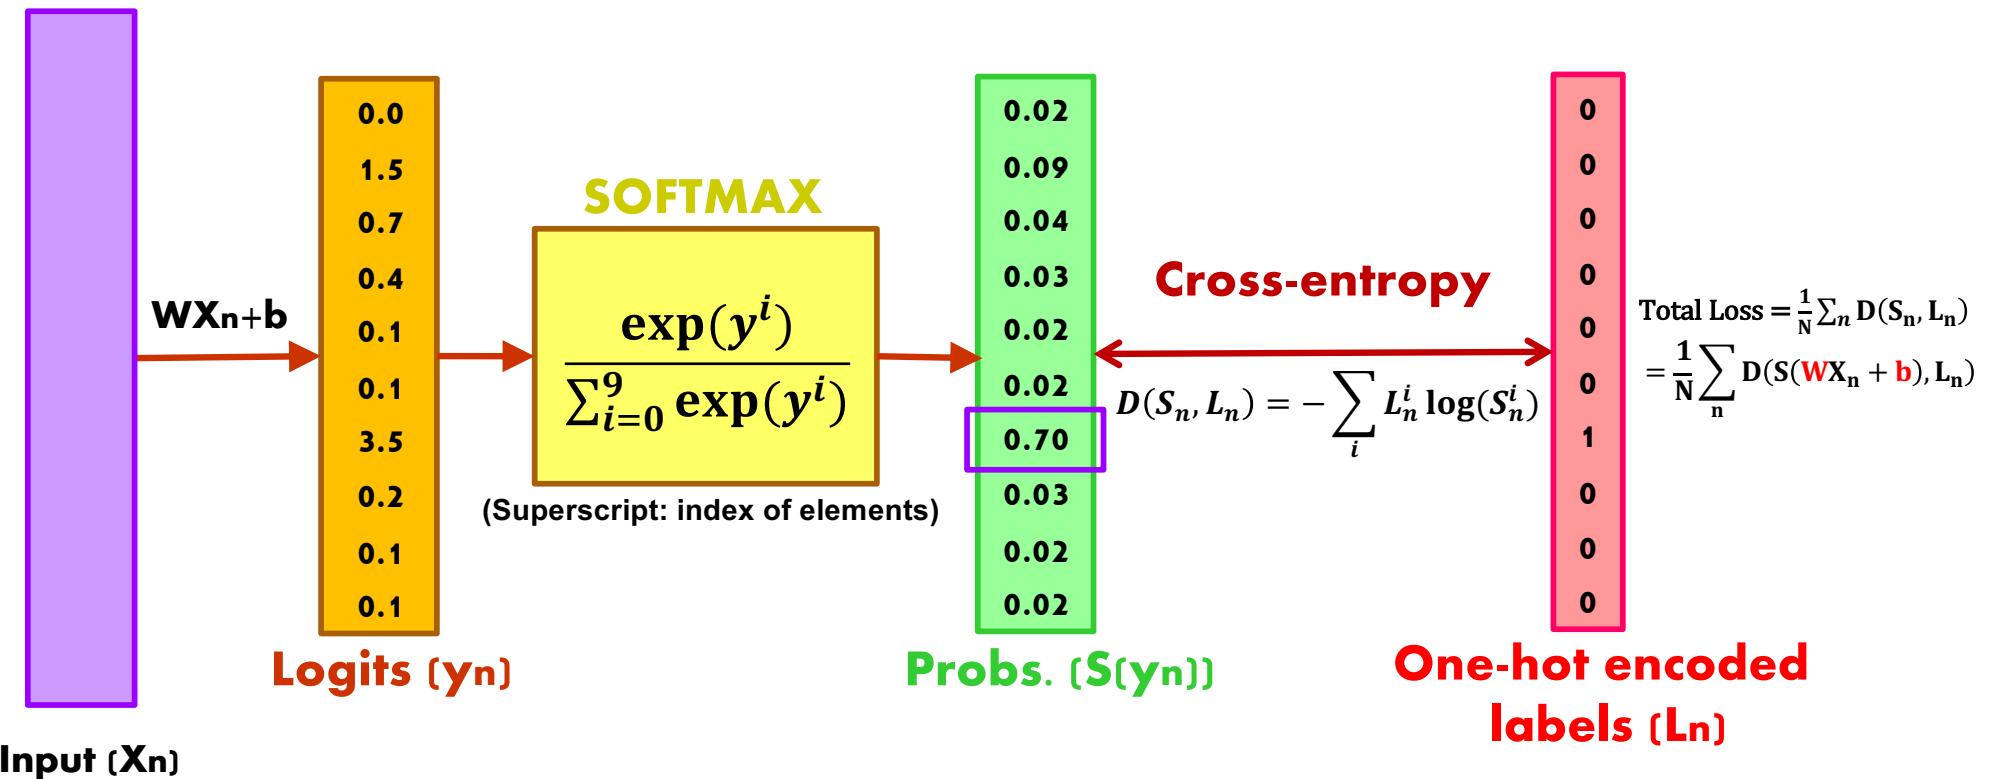
model.fit(x_train, y_train, epochs=5, batch_size=128)
model.evaluate(x_test, y_test)
```

Logistic (i.e., linear) classifier

Set of N labeled inputs: $D = \{(X_1, y_1), \dots, (X_N, y_N)\}$ $\begin{cases} X = \text{input data} \\ y = \text{input label} \end{cases}$



Logistic Classifier (linear classifier)



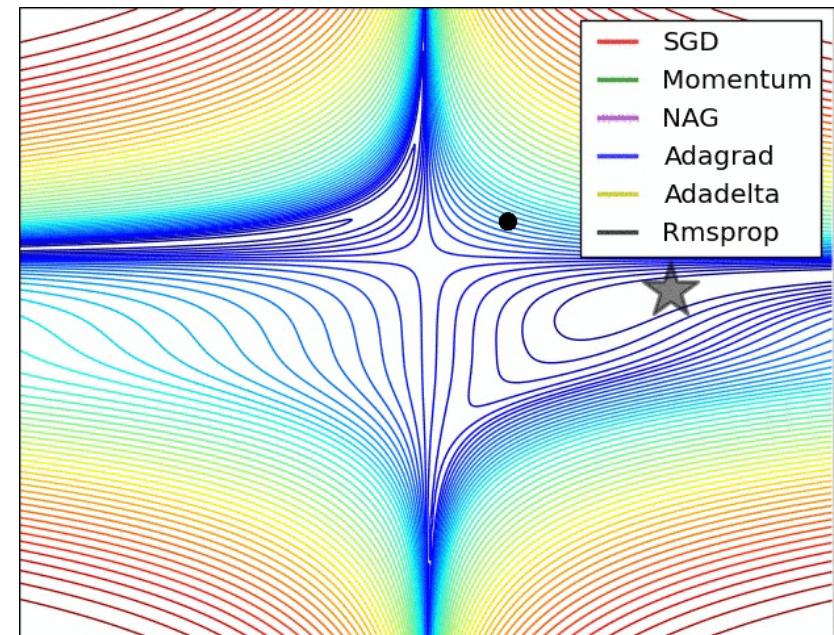
Optimizer

$$L = \frac{1}{N} \sum_n D(S(WX_n + b), L_n)$$

Gradient descent:

$$w := w - \alpha \frac{\partial L}{\partial w}$$

Learning rate



[An overview of gradient descent optimization algorithms](#)

Logistic Classifier (linear classifier)

- Disadvantages:
 - Few parameters: $\#parameters = 10 \times 784 + 10 = 7850$
 - At the end, it's a linear model!
- Advantages:
 - Linear models are STABLE!